



Inspire...Educate...Transform.

Statistics and Probability in Decision Modeling

**Principal Components Analysis (PCA),
Regularization (Ridge, LASSO and Elastic Nets
Regression), Time Series Forecasting (Part 1)**

**Dr. Sridhar Pappu
Executive VP – Academics, INSOFE**

January 20, 2018

Slides on PCA and Regularization courtesy Dr. Anand Jayaraman



cmcott. 04/13/12 #135

CSE 7302c



VIF, GVIF, GVIF $\left(\frac{1}{2*df}\right)$

Predicting Coronary Heart Disease Case

```
call:  
glm(formula = TenYearCHD ~ ., family = binomial, data = train)  
  
Deviance Residuals:  
    Min      1Q      Median      3Q      Max  
-1.9392 -0.5998 -0.4211 -0.2771  2.8632  
  
Coefficients:  
              Estimate Std. Error z value Pr(>|z|)  
(Intercept) -8.360272  0.864696 -9.668 < 2e-16 ***  
male          0.524080  0.130836  4.006 6.19e-05 ***  
age           0.065429  0.008049  8.129 4.34e-16 ***  
education     -0.041105  0.059185 -0.695 0.487366  
currentSmoker  0.120498  0.187629  0.642 0.520735  
cigsPerDay    0.016471  0.007488  2.200 0.027825 *  
BPMeds         0.169118  0.282140  0.599 0.548898  
prevalentstroke 1.156666  0.560179  2.065 0.038940 *  
prevalentHyp   0.307077  0.166034  1.849 0.064389 .  
diabetes       -0.319937  0.392574 -0.815 0.415087  
totChol        0.003799  0.001330  2.856 0.004290 **  
sysBP          0.011144  0.004446  2.507 0.012188 *  
diaBP          -0.001861  0.007760 -0.240 0.810517  
BMI            0.008812  0.015662  0.563 0.573702  
heartRate      -0.007273  0.005131 -1.418 0.156296  
glucose         0.009227  0.002752  3.353 0.000798 ***  
---  
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
  
(Dispersion parameter for binomial family taken to be 1)  
  
Null deviance: 2176.6 on 2565 degrees of freedom  
Residual deviance: 1919.9 on 2550 degrees of freedom  
(402 observations deleted due to missingness)  
AIC: 1951.9
```

- Education as **numeric** variable

```
> framingham = read.csv("framingham.csv")  
> str(framingham)  
'data.frame': 4240 obs. of 16 variables:  
 $ male      : int 1 0 1 0 0 0 0 1 1 ...  
 $ age       : int 39 46 48 61 46 43 63 45 52 43 ...  
 $ education : int 4 2 1 3 3 2 1 2 1 1 ...
```

- car package gives the following VIF figures

```
> car::vif(framinghamLog)  
               male             age            education  
1.247670      1.278996      1.057810
```

CSE 7302C



VIF, GVIF, GVIF $\left(\frac{1}{2*df}\right)$

Predicting Coronary Heart Disease Case

```

Call:
glm(formula = TenYearCHD ~ ., family = binomial, data = train)

Deviance Residuals:
    Min      1Q  Median      3Q     Max 
-1.9315 -0.5948 -0.4196 -0.2733  2.8925 

Coefficients:
            Estimate Std. Error z value Pr(>|z|)    
(Intercept) -8.286327  0.856000 -9.680 < 2e-16 ***
male         0.506116  0.131991  3.834 0.000126 *** 
age          0.064020  0.008135  7.869 3.56e-15 ***
education2  -0.210890  0.148310 -1.422 0.155038  
education3  -0.120464  0.174595 -0.690 0.490220  
education4  -0.082216  0.199469 -0.412 0.680212  
currentSmoker 0.125147  0.187693  0.667 0.504921  
cigsPerDay   0.016589  0.007482  2.217 0.026611 *  
BPMeds       0.177341  0.282452  0.628 0.530094  
prevalentStroke 1.158996  0.563826  2.056 0.039822 *  
prevalentHyp   0.308709  0.166304  1.856 0.063412 .  
diabetes      -0.318608  0.393231 -0.810 0.417807  
totChol       0.003860  0.001334  2.894 0.003801 ** 
sysBP         0.011195  0.004451  2.515 0.011893 *  
diaBP         -0.001726  0.007766 -0.222 0.824120  
BMI           0.007535  0.015680  0.481 0.630828  
heartRate     -0.007132  0.005136 -1.389 0.164946  
glucose        0.009221  0.002748  3.356 0.000791 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 2176.6 on 2565 degrees of freedom
Residual deviance: 1918.2 on 2548 degrees of freedom
(402 observations deleted due to missingness)
AIC: 1954.2

Number of Fisher Scoring iterations: 5

```

- Education as **categorical** variable


```

> framingham$education = factor(framingham$education)
> str(framingham)
'data.frame': 4240 obs. of 16 variables:
 $ male      : int 1 0 1 0 0 0 0 1 1 ...
 $ age       : int 39 46 48 61 46 43 63 45 52 43 ...
 $ education : Factor w/ 4 levels "1","2","3","4": 4 2
 1 3 3 2 1 2 1 1 ...
      
```
- car package gives the following (G)VIF figures


```

> car::vif(framinghamLog)
              GVIF Df GVIF^(1/(2*Df))
male          1.268872 1    1.126442
age          1.302896 1    1.141445
education    1.121533 3    1.019300
      
```
-  Use the **square** of the GVIF $\left(\frac{1}{2*df}\right)$ value and apply the VIF rule of thumb

CSEZ302C





Present Day



Dealing with a large number of dimensions

DIMENSIONALITY REDUCTION



Understanding the concept

PRINCIPAL COMPONENTS ANALYSIS (PCA)

PCA

- Linear transformation technique used for dimensionality reduction
- Converts a set of observations of possibly correlated variables into a set of values of linearly uncorrelated variables (principal components)
- Explains a large part of the variation in data using a small number of principal components

CSE 7302c



PCA Applications

- Commonly used in stock market predictions, analysis of gene expression (process of using information from a gene to synthesize proteins, etc.) data, etc.
- Many applications in taxonomy, finance, biology, neuroscience, pharmacy, agriculture, ecology, health, architecture, etc.

CSE 7302c



Algorithmic Trading

- Analyze patterns in economic data and past stock prices to forecast future price movements
- Target Variable: Stock Price Direction (Up or Down), e.g., Bank Nifty

CSE 7302c



Algorithmic Trading

- Predictor Variables: 352 different variables quantifying economic data and price movements
 - Past data of Bank Nifty and quantifying patterns therein
 - Quantifying patterns in the US market on the previous data
 - Patterns in the variations in Dollar-Rupee exchange rate
 - Interest rate changes in India
 - Interest rate changes in US
- The patterns, not values, over the past week, month, etc. are studied after downloading all this data

CSE 7302c



Algorithmic Trading

- Many of the predictors may be irrelevant
- Many of the predictors may be similar (correlated)
- Correlation could be complex, e.g., 5 variables put together in some manner may be correlated to another variable
- Approximately 14 hours to do the simulation

CSE 7302c



Algorithmic Trading

- Is there a way to figure out which variables are important and which are not essential?

CSE 7302c



Which is the relevant variable?

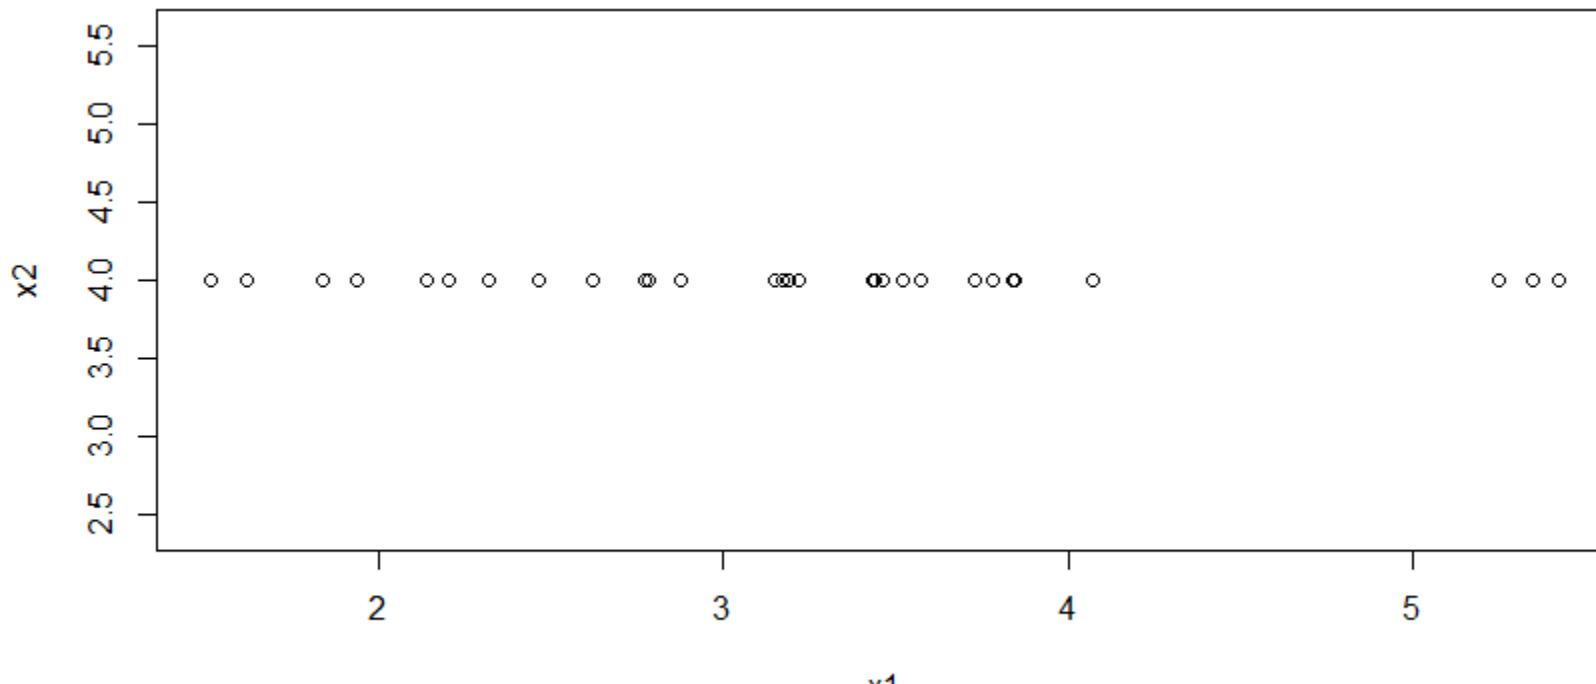
y	x1	x2
21	2.62	4
21	2.875	4
22.8	2.32	4
21.4	3.215	4
18.7	3.44	4
18.1	3.46	4
14.3	3.57	4
24.4	3.19	4
22.8	3.15	4
19.2	3.44	4
17.8	3.44	4
16.4	4.07	4
17.3	3.73	4
15.2	3.78	4
10.4	5.25	4
10.4	5.424	4
14.7	5.345	4
32.4	2.2	4

There is no variation in x2. All the variation is only along x1. So y cannot be explained by x2.

$$sd(x_1) = 0.98$$

$$sd(x_2) = 0$$

We didn't even need to look at *y* to make the conclusion!



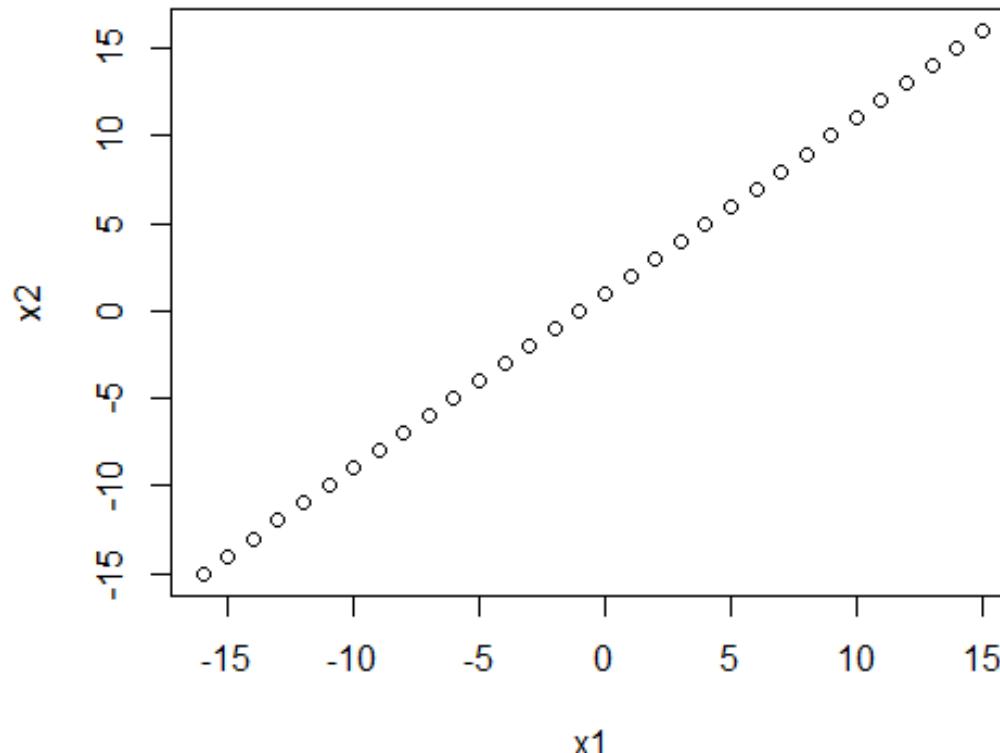
All the variation in the explanatory variables is in x_1 direction only.

This direction is called the dominant Principal Component of the explanatory variables.

The x_2 direction is redundant, since there is no variation along that axis.

How many relevant features are here?

y	x1	x2
21	-16	-15
21	-15	-14
22.8	-14	-13
21.4	-13	-12
18.7	-12	-11
18.1	-11	-10
14.3	-10	-9
24.4	-9	-8
22.8	-8	-7
19.2	-7	-6
17.8	-6	-5
16.4	-5	-4
17.3	-4	-3
15.2	-3	-2
10.4	-2	-1
10.4	-1	0
14.7	0	1
32.4	1	2
30.4	2	3
33.9	3	4



This seems to have variation in both x_1 and x_2 . However, the data points in x_1 and x_2 are perfectly correlated.

How many relevant features are here?

y	x1	x2
21	-16	-15
21	-15	-14
22.8	-14	-13
21.4	-13	-12
18.7	-12	-11
18.1	-11	-10
14.3	-10	-9
24.4	-9	-8
22.8	-8	-7
19.2	-7	-6
17.8	-6	-5
16.4	-5	-4
17.3	-4	-3
15.2	-3	-2
10.4	-2	-1
10.4	-1	0
14.7	0	1
32.4	1	2
30.4	2	3
33.9	3	4

If we create new variables

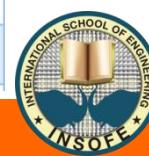
$$X_1 = x_2 + x_1$$

$$X_2 = x_2 - x_1$$

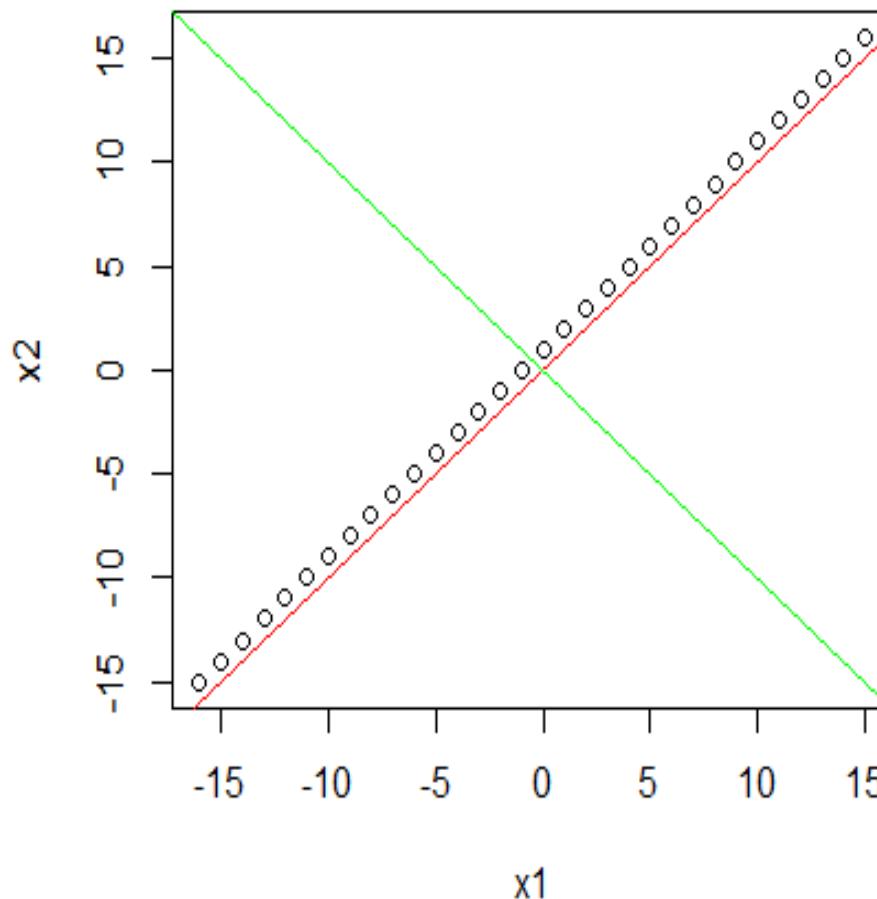
In terms of new variables X_1 and X_2 it is clear that there is only one true relevant feature.

y	X1	X2
21	-31	1
21	-29	1
22.8	-27	1
21.4	-25	1
18.7	-23	1
18.1	-21	1
14.3	-19	1
24.4	-17	1
22.8	-15	1
19.2	-13	1
17.8	-11	1
16.4	-9	1
17.3	-7	1
15.2	-5	1
10.4	-3	1
10.4	-1	1
14.7	1	1
32.4	3	1
30.4	5	1

CSE 7302C



Transformed Variables: Interpretation



Transformed Variables: Interpretation

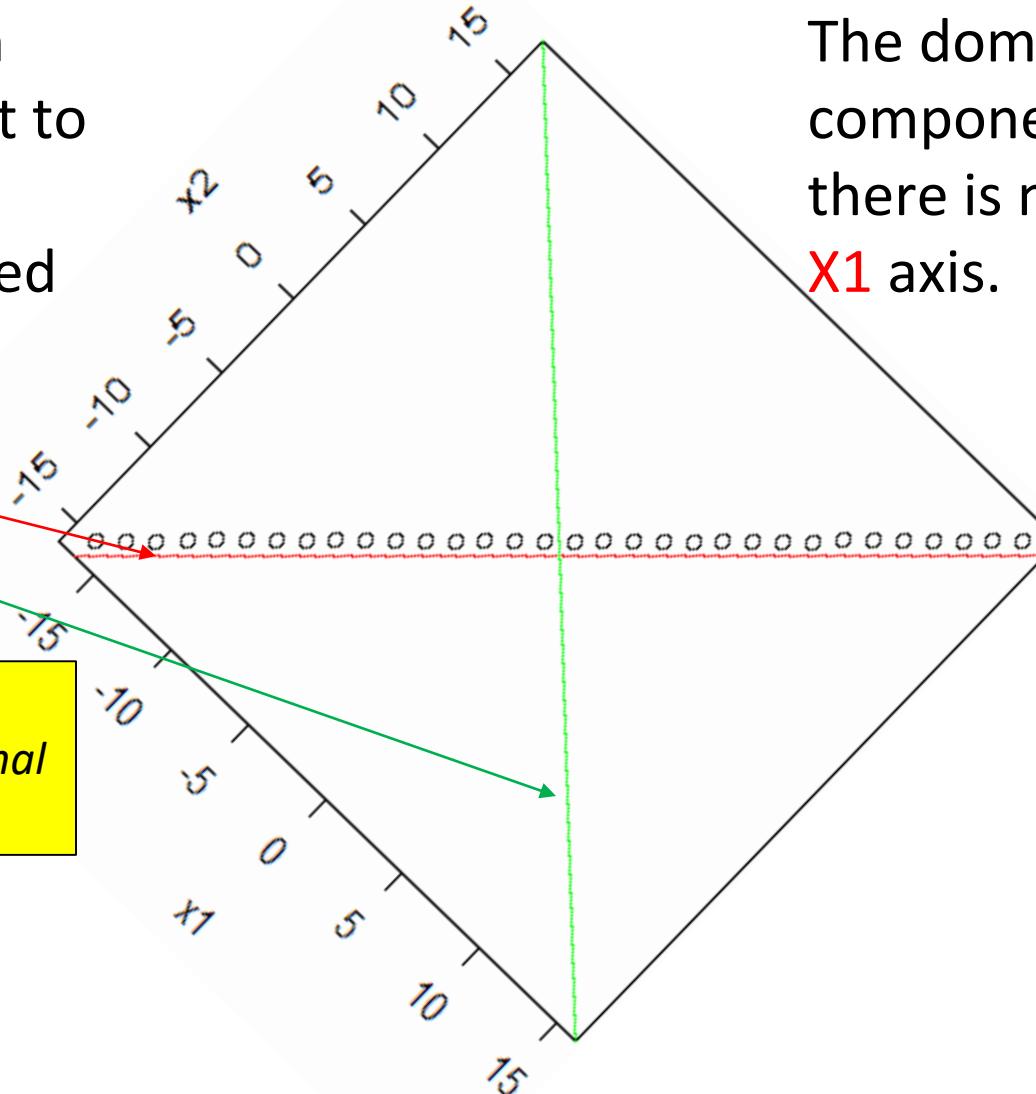
The transformation we did is equivalent to viewing the data points from a rotated coordinate system.

The dominant Principal component (along which there is maximum variation) is **X1** axis.

Red = X1 axis

Green = X2 axis

Note both **X1** and **X2** are combinations of the original variables, x_1 and x_2 .



Standardization is a Prerequisite

y	x1	x2
21	-16	-15
21	-15	-14
22.8	-14	-13
21.4	-13	-12
18.7	-12	-11
18.1	-11	-10
14.3	-10	-9
24.4	-9	-8
22.8	-8	-7
19.2	-7	-6
17.8	-6	-5
16.4	-5	-4
17.3	-4	-3
15.2	-3	-2
10.4	-2	-1
10.4	-1	0
14.7	0	1
32.4	1	2
30.4	2	3
33.9	3	4

Suppose **y** is miles per gallon, **x1** is weight of a car and **x2** is horsepower. Does it make sense to transform variables (x_1+x_2 or x_1-x_2) with different scales and units?

y	X1	X2
21	-31	1
21	-29	1
22.8	-27	1
21.4	-25	1
18.7	-23	1
18.1	-21	1
14.3	-19	1
24.4	-17	1
22.8	-15	1
19.2	-13	1
17.8	-11	1
16.4	-9	1
17.3	-7	1
15.2	-5	1
10.4	-3	1
10.4	-1	1
14.7	1	1
32.4	3	1
30.4	5	1

No. Before running PCA, it is a **must** to **standardize all variables**.

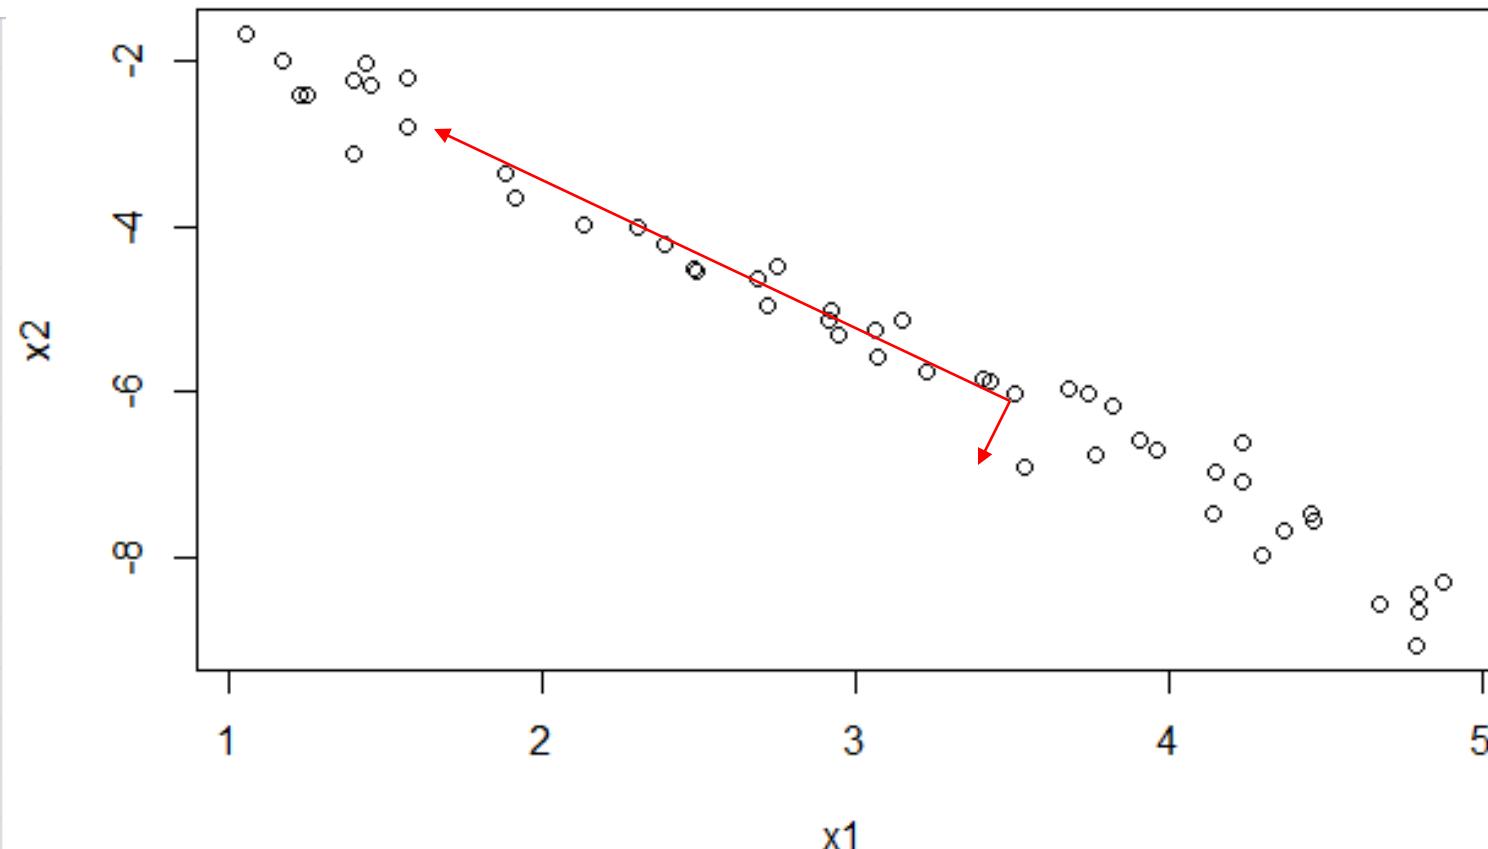
Principal Components Analysis

PCA is the method which allows you to identify the “directions” in which most of the variation in the data is present. Note we do not consider y at all in this analysis.

Equivalently, it can be thought of as a method to identify the “directions” along which there is least variation (or least useful info). Identifying this would allow us to drop these irrelevant directions from our regressions/model building.

Principal Component Directions

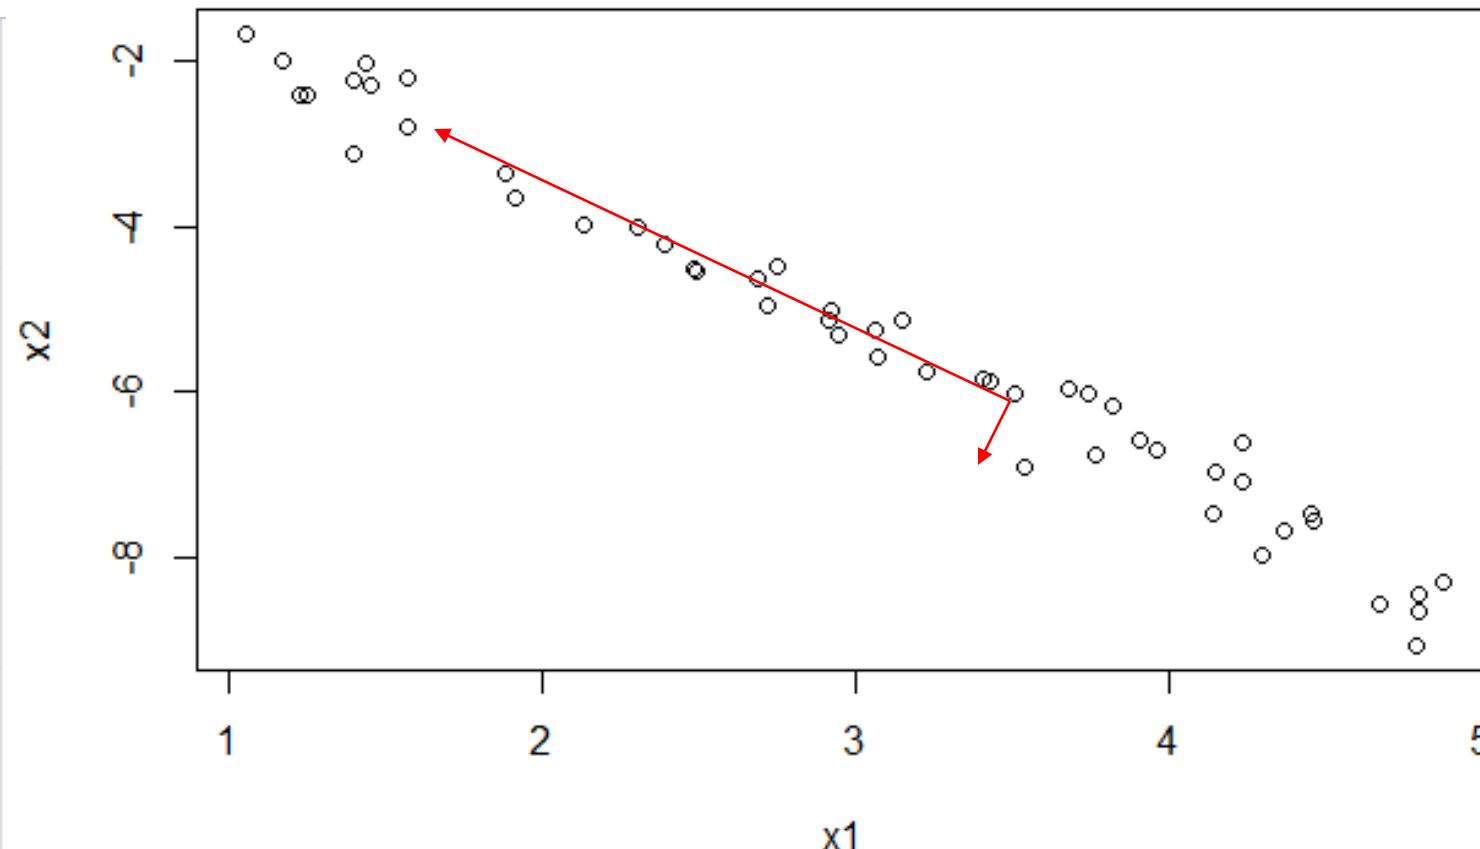
y	x1	x2
-42.4957832	4.797034	-8.651047
-2.0248387	1.220774	-2.428135
-1.9379560	1.437964	-2.037009
-13.0806296	2.719718	-4.960111
-35.1968731	4.302478	-7.967436
-31.7650916	4.139152	-7.471322
-22.4288444	3.681965	-5.951819
-21.4937406	3.512183	-6.016649
-0.9004686	1.050161	-1.683525
-3.3677489	1.398850	-3.111843
-26.3779407	3.907309	-6.573550
-34.6572233	4.460774	-7.552211
-24.1480305	3.819653	-6.157824
-23.1244874	3.744704	-6.023781
-2.3255245	1.452426	-2.284816
-29.7199029	4.153706	-6.950889
-2.1284295	1.393772	-2.233364
-16.0110878	3.060388	-5.258223



Can you identify the Principal Components?

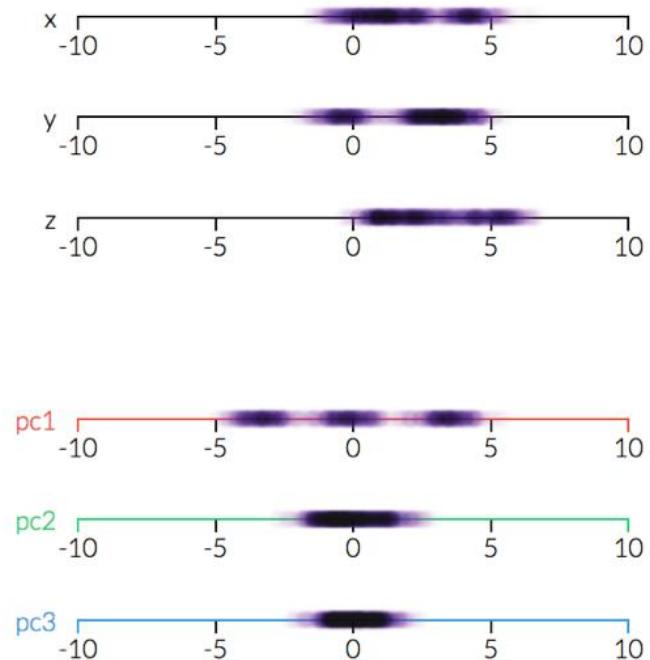
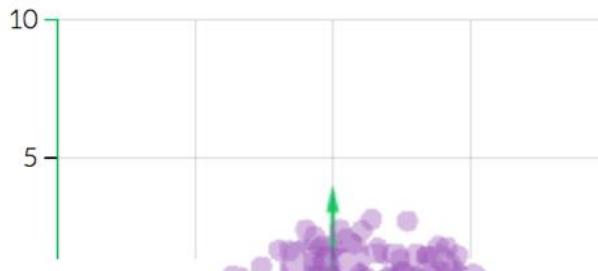
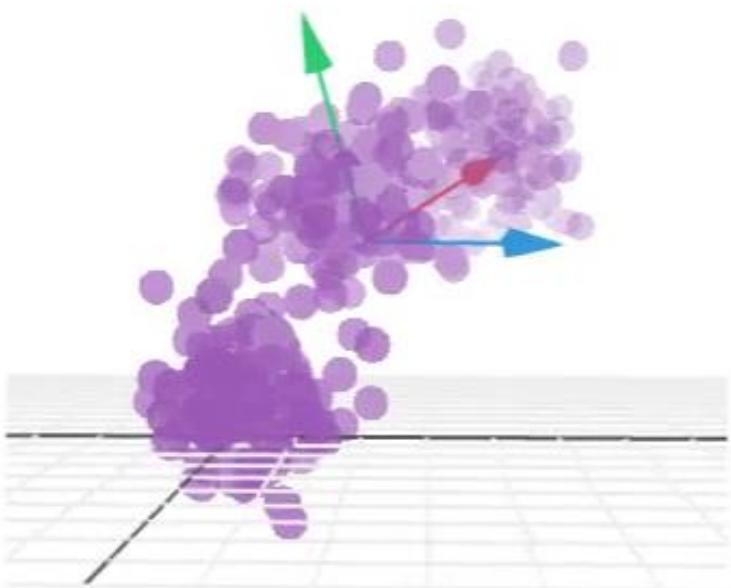
Dimension Reduction

y	x1	x2
-42.4957832	4.797034	-8.651047
-2.0248387	1.220774	-2.428135
-1.9379560	1.437964	-2.037009
-13.0806296	2.719718	-4.960111
-35.1968731	4.302478	-7.967436
-31.7650916	4.139152	-7.471322
-22.4288444	3.681965	-5.951819
-21.4937406	3.512183	-6.016649
-0.9004686	1.050161	-1.683525
-3.3677489	1.398850	-3.111843
-26.3779407	3.907309	-6.573550
-34.6572233	4.460774	-7.552211
-24.1480305	3.819653	-6.157824
-23.1244874	3.744704	-6.023781
-2.3255245	1.452426	-2.284816
-29.7199029	4.153706	-6.950889
-2.1284295	1.393772	-2.233364
-16.0110878	3.060388	-5.258223



If we had to simplify and pick only one feature, it is easier to do so in the transformed variables. We can keep the dominant principal component and skip the non-dominant component.

An example in 3D



Source: <http://setosa.io/ev/principal-component-analysis/>
Last accessed: January 19, 2018

CSE 7302c



Principal Components Analysis – Summarizing the Concept

- PCA is a way for identifying the directions of largest to smallest variance
- It also tells how to transform the data into the new coordinate system
- Once we transform the data into the new coordinate system, we can choose to drop the variables which do not have much variance
- This gives a way to **reduce dimensions** and focus on ones with largest variance
- PCA is a **dimensionality reduction** technique, not a **feature selection** method, i.e., we don't select variables but we combine them to find fewer dimensions explaining most variance in the data

CSE
7302c



Array, Vector, Matrix...THE MATH METHODOLOGY

CSE 7302c



Array! Matrix! And other **Nightmares!**



INTERNAL - II

Computer programming

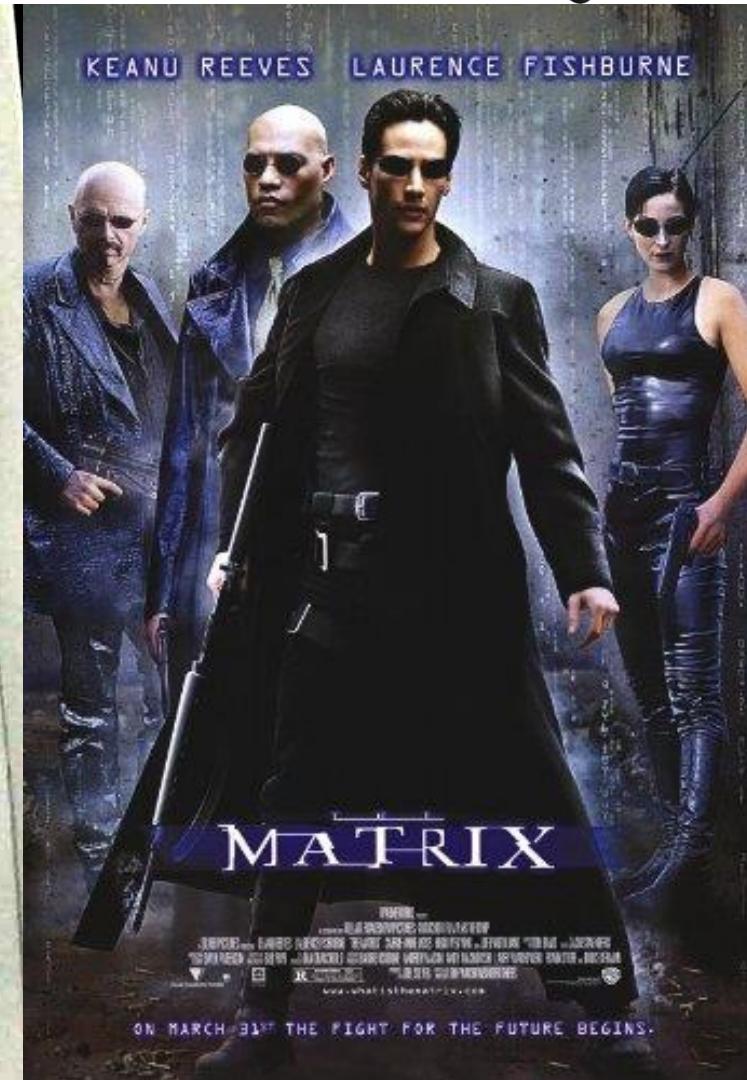
1. Define Array.

Ans: * **Array:**

An Array is used to call a boy or a person who is at a distance far away from us who are visible to our naked eye.

Example: Array Rupesh.

Meet me
here



Array! Matrix! They are SIMPLE!



What is an array?

Dimensions

Example

1

0	1	2
---	---	---

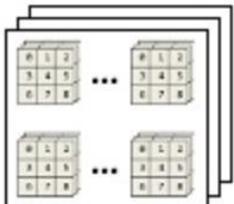
2

0	1	2
3	4	5
6	7	8

3

0	1	2
3	4	5
6	7	8

N



Terminology

Vector

(1D Array)

Matrix

(2D Array)

3D Array
(3rd order Tensor)

ND Array

WineClass	Alcohol	MalicAcid	Ash	AlkalinityAsh
barolo	14.23	1.71	2.43	15.6
barolo	13.2	1.78	2.14	11.2
barolo	13.16	2.36	2.67	18.6
barolo	14.37	1.95	2.5	16.8
barolo	13.24	2.59	2.87	21
barolo	14.2	1.76	2.45	15.2
barolo	14.39	1.87	2.45	14.6
barolo	14.06	2.15	2.61	17.6
barolo	14.83	1.64	2.17	14
barolo	13.86	1.35	2.27	16
barolo	14.1	2.16	2.3	18
barolo	14.12	1.48	2.32	16.8
barolo	13.75	1.73	2.41	16
barolo	14.75	1.73	2.39	11.4
barolo	14.38	1.87	2.38	12

Source: <http://www.java67.com/2014/08/what-is-array-data-structure-in-java.html>
Last accessed: July 07, 2017

CSE 7302C



Regression in Matrix Form

Market Price (\$1000) (y)	Area (sq ft) (x1)	Age of House (years) (x2)
63	1605	35
65.1	2489	45
69.9	1553	20
76.8	2404	32
73.9	1884	25
77.9	1558	14
74.9	1748	8
78	3105	10
79	1682	28
83.4	2470	30
79.5	1820	2
83.9	2143	6
79.7	2121	14
84.5	2485	9
96	2300	19
109.5	2714	4
102.5	2463	5
121	3076	7
104.9	3048	3
128	3267	6
129	3069	10
117.9	4765	11
140	4540	8

	Coefficients	Standard Error	t Stat	P-value
Intercept	57.35074586	10.00715186	5.73097587	1.31298E-05
Area (sq ft) (x1)	0.017718036	0.00314562	5.632605205	1.63535E-05
Age of House (years) (x2)	-0.666347946	0.227996703	-2.922620973	0.008417613

$$\hat{Y} = X\beta$$

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \hat{y}_3 \\ \hat{y}_4 \end{bmatrix} = \begin{bmatrix} 1 & 1605 & 35 \\ 1 & 2489 & 45 \\ 1 & 1553 & 20 \\ 1 & 2404 & 32 \end{bmatrix} \begin{bmatrix} 57.3507 \\ 0.0177 \\ -0.6663 \end{bmatrix}$$

$$Y = X\beta + \varepsilon$$

PCA Methodology

Start with analyzing the covariance matrix of the features.

x1	x2
4.797034	-8.651047
1.220774	-2.428135
1.437964	-2.037009
2.719718	-4.960111
4.302478	-7.967436
4.139152	-7.471322
3.681965	-5.951819
3.512183	-6.016649
1.050161	-1.683525
1.398850	-3.111843
3.907309	-6.573550
4.460774	-7.552211
3.819653	-6.157824
3.744704	-6.023781
1.452426	-2.284816
4.153706	-6.950889

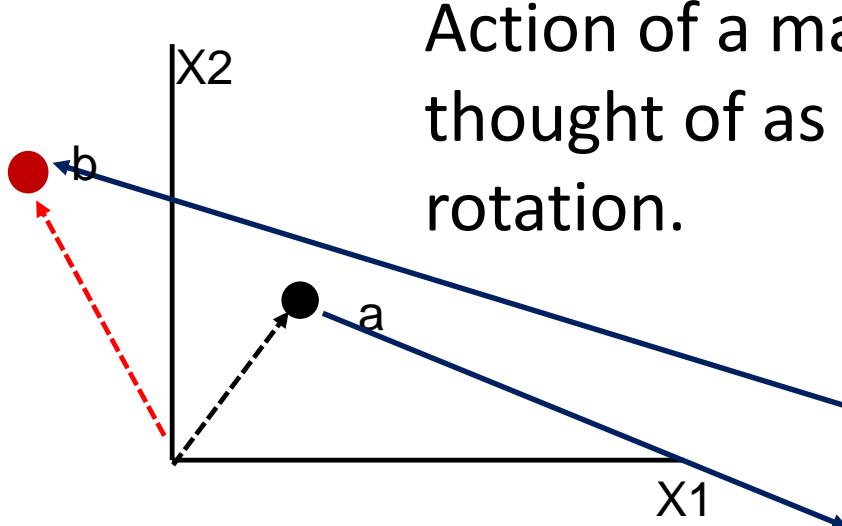
$$C = \begin{bmatrix} cov(x_1, x_1) & cov(x_1, x_2) \\ cov(x_2, x_1) & cov(x_2, x_2) \end{bmatrix}$$

```
> cov(d[,c("x1", "x2")])
```

	x1	x2
x1	1.358045	-2.35219
x2	-2.352190	4.18192

The covariance matrix contains information about both correlations and the “special directions” of maximal variances.

Matrix as a Transformation on a Vector

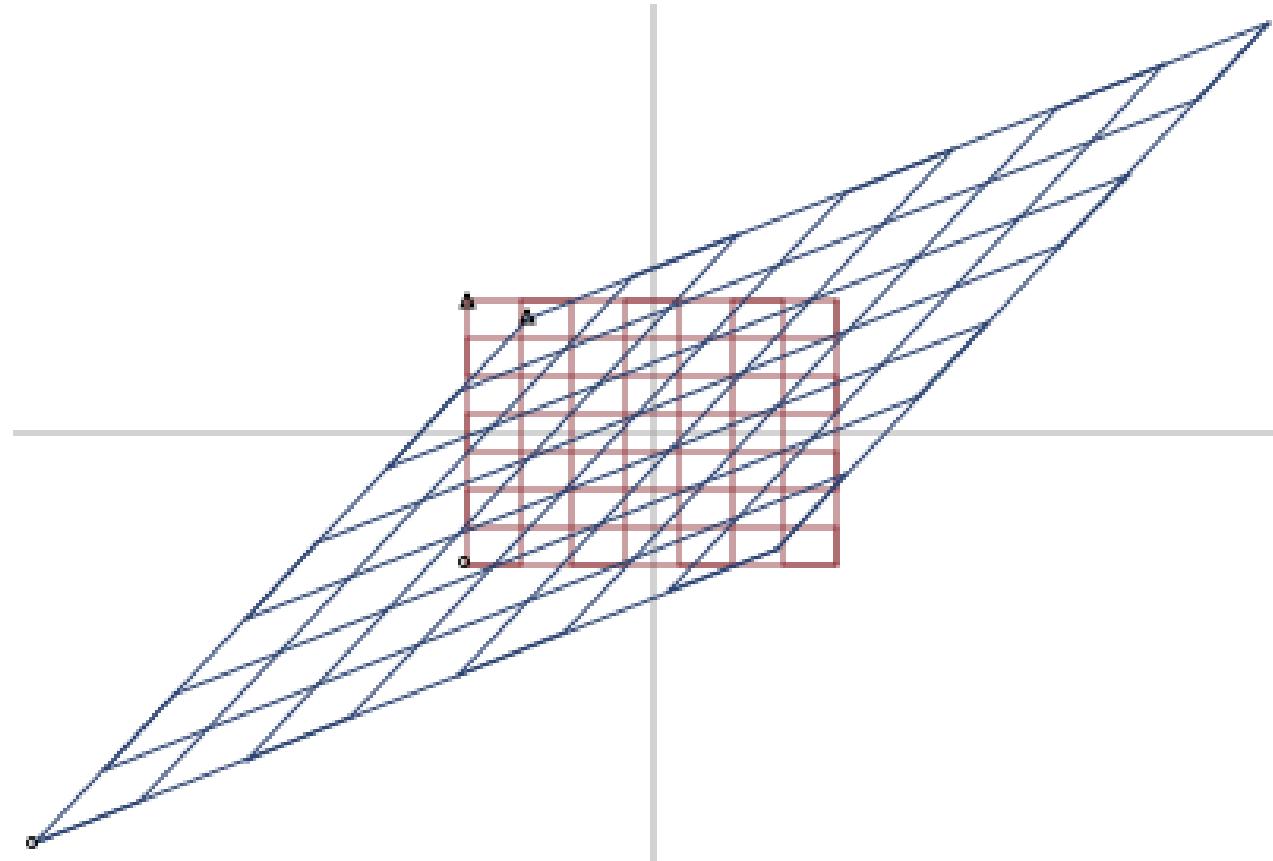


Action of a matrix on a general vector can be thought of as a combination of stretch and rotation.

$$\begin{bmatrix} 1.36 & -2.35 \\ -2.35 & 4.18 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} -0.99 \\ 1.83 \end{bmatrix}$$

```
> cov(d[,c("x1","x2")])  
           x1          x2  
x1  1.358045 -2.35219  
x2 -2.352190  4.18192
```

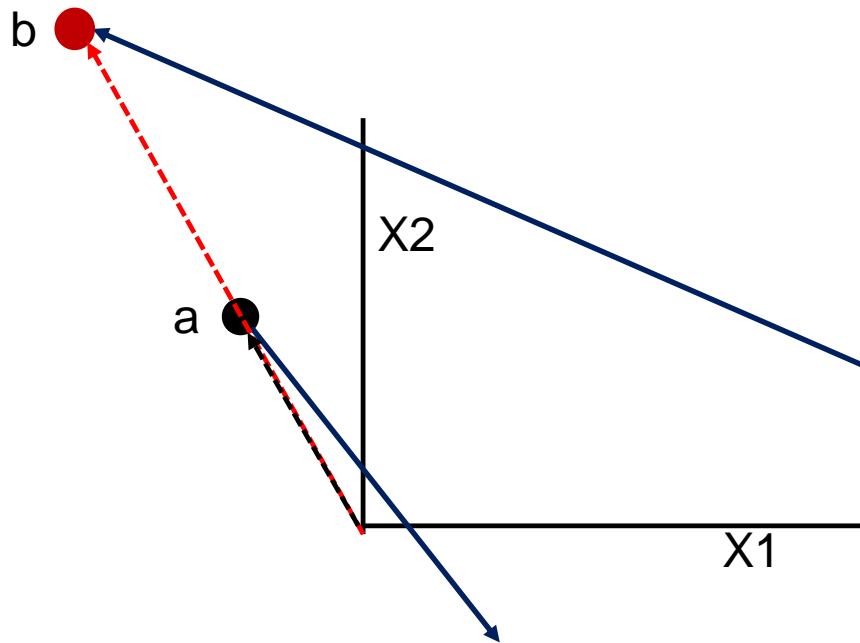
A Transformation Matrix that Stretches, Rotates and Skews



CSE 7302c



Matrix as a Transformation on a Vector



$$\begin{bmatrix} 1.36 & -2.35 \\ -2.35 & 4.18 \end{bmatrix} \begin{bmatrix} -0.49 \\ 0.87 \end{bmatrix} = \begin{bmatrix} -2.715 \\ 4.795 \end{bmatrix} = 5.51 \begin{bmatrix} -0.49 \\ 0.87 \end{bmatrix}$$

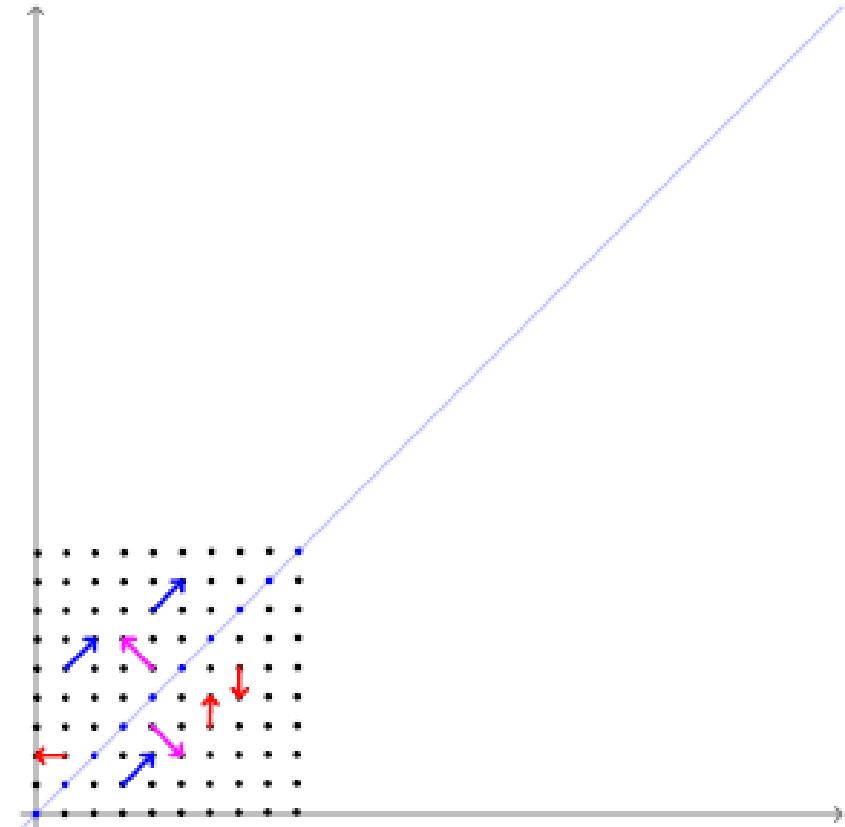
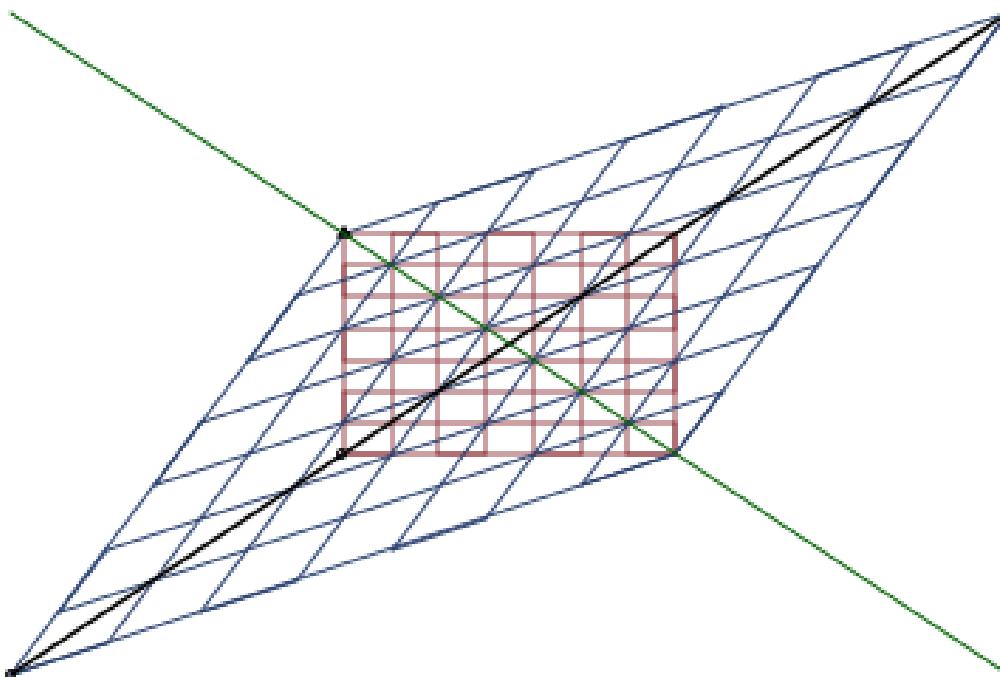
CS
E 7302c



Eigenvectors

Eigen in German means “very own”. For example, “mein eigenes auto” means “my very own car”.

So, an eigenvector describes the relationship with itself and not with anything else.



CSE 7302c

Source: <https://deeplearning4j.org/eigenvector>
Last accessed: July 07, 2017



Eigenvectors Mathematics

- The eigenvectors and eigenvalues of matrix A are defined to be the nonzero \mathbf{x} and λ values, respectively, that solve

$$\mathbf{Ax} = \lambda \mathbf{x} \quad (\mathbf{A} \text{ is just stretching})$$
$$\begin{bmatrix} 1.36 & -2.35 \\ -2.35 & 4.18 \end{bmatrix} \begin{bmatrix} -0.49 \\ 0.87 \end{bmatrix} = 5.51 \begin{bmatrix} -0.49 \\ 0.87 \end{bmatrix}$$

- For a n -dim square matrix, there are at most n eigenvectors and n eigenvalues.



Eigenvectors and PCA

- Eigenvectors are the Principal Component directions
- Eigenvalues are the magnitude of stretch
- Eigenvalues represent the variance along those (new) directions.

CSE 7302c



```
> c <- cov(d[,c("x1","x2")])
```

```
> c
```

	x1	x2
x1	1.358045	-2.35219
x2	-2.352190	4.18192

```
> eigen(c)
```

```
$values
```

```
[1] 5.51340440 0.02656068
```

```
$vectors
```

```
[,1]
```

```
[,2]
```

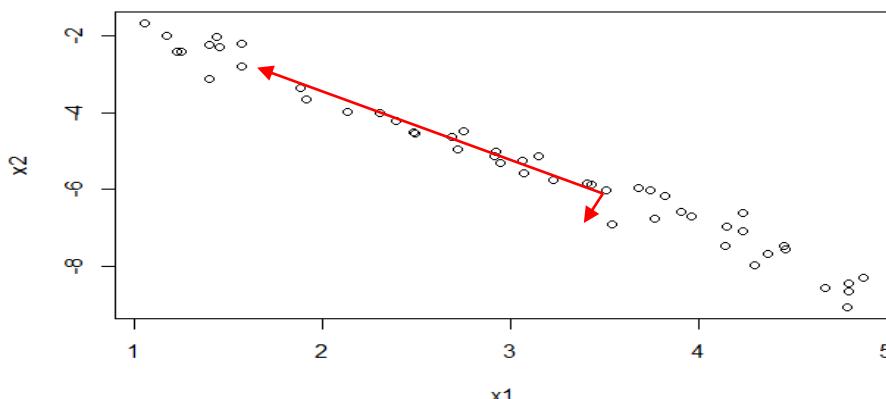
```
[1,] -0.4926140 -0.8702479
```

```
[2,] 0.8702479 -0.4926140
```

$$E_1 = \begin{bmatrix} -0.49 \\ 0.87 \end{bmatrix} \text{ with } \lambda_1 = 5.51$$

$$E_2 = \begin{bmatrix} -0.87 \\ -0.49 \end{bmatrix} \text{ with } \lambda_2 = 0.02$$

Dominant Principal Component



The relative size of eigenvalues says variance in one direction is much higher than in the other.

CSE 7302c



```

> c <- cov(d[,c("x1","x2")])
> c
      x1      x2
x1  1.358045 -2.35219
x2 -2.352190  4.18192
> eigen(c)
$values
[1] 5.51340440 0.02656068

```

```

$vectors
      [,1]      [,2]
[1,] -0.4926140 -0.8702479
[2,]  0.8702479 -0.4926140

```

Remember, in the other example, when we transformed the data points from original variables x_1 , x_2 into new transformed variables, X_1 (x_2+x_1) and X_2 (x_2-x_1), we could reduce the dimensions?

This matrix of eigenvectors as a whole also allows you to transform each one of our data-points into new variables X_1 and X_2 .

Transform into Principal Components

y	x1	x2
-42.4957832	4.797034	-8.651047
-2.0248387	1.220774	-2.428135
-1.9379560	1.437964	-2.037009
-13.0806296	2.719718	-4.960111
-35.1968731	4.302478	-7.967436
-31.7650916	4.139152	-7.471322
-22.4288444	3.681965	-5.951819
-21.4937406	3.512183	-6.016649
-0.9004686	1.050161	-1.683525
-3.3677489	1.398850	-3.111843
-26.3779407	3.907309	-6.573550
-34.6572233	4.460774	-7.552211
-24.1480305	3.819653	-6.157824
-23.1244874	3.744704	-6.023781
-2.3255245	1.452426	-2.284816
-29.7199029	4.153706	-6.950889
-2.1284295	1.393772	-2.233364
-16.0110878	3.060388	-5.258223

Any record in our data set (x_1, x_2) when multiplied by the matrix of eigenvectors, we get the new coordinates in the rotated principal component axis.

$$\begin{bmatrix} X_1 \\ X_2 \end{bmatrix} = \begin{bmatrix} -0.49 & -0.87 \\ 0.87 & -0.49 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}$$

$$\begin{bmatrix} X_1 \\ X_2 \end{bmatrix} = \begin{bmatrix} -0.49x_1 - 0.87x_2 \\ 0.87x_1 - 0.49x_2 \end{bmatrix}$$

x1	x2
4.797034	-8.651047
1.220774	-2.428135
1.437964	-2.037009
2.719718	-4.960111
4.302478	-7.967436
4.139152	-7.471322
3.681965	-5.951819
3.512183	-6.016649
1.050161	-1.683525
1.39885	-3.111843
3.907309	-6.57355
4.460774	-7.552211
3.819653	-6.157824
3.744704	-6.023781
1.452426	-2.284816
4.153706	-6.950889
1.393772	-2.233364
3.060388	-5.258223
4.788337	-9.048708
2.918199	-5.015753



X1	X2
-9.891641591	0.087018102
-2.714449748	0.133757285
-2.481064002	-0.247926
-5.656297344	0.076591242
-9.053105345	0.180648073
-8.540906504	0.07838948
-6.993345492	-0.272272944
-6.966126674	-0.09258435
-1.982408107	-0.084572421
-3.39716793	0.315591153
-7.645413199	-0.162104692
-8.769735486	-0.161654336
-7.240447948	-0.29061469
-7.086876382	-0.291421939
-2.703841707	-0.138438327
-8.095170283	-0.19064869
-2.630171931	-0.112740783
-6.083547498	-0.073021965
-10.23342098	0.290480024
-5.802494197	-0.068726403

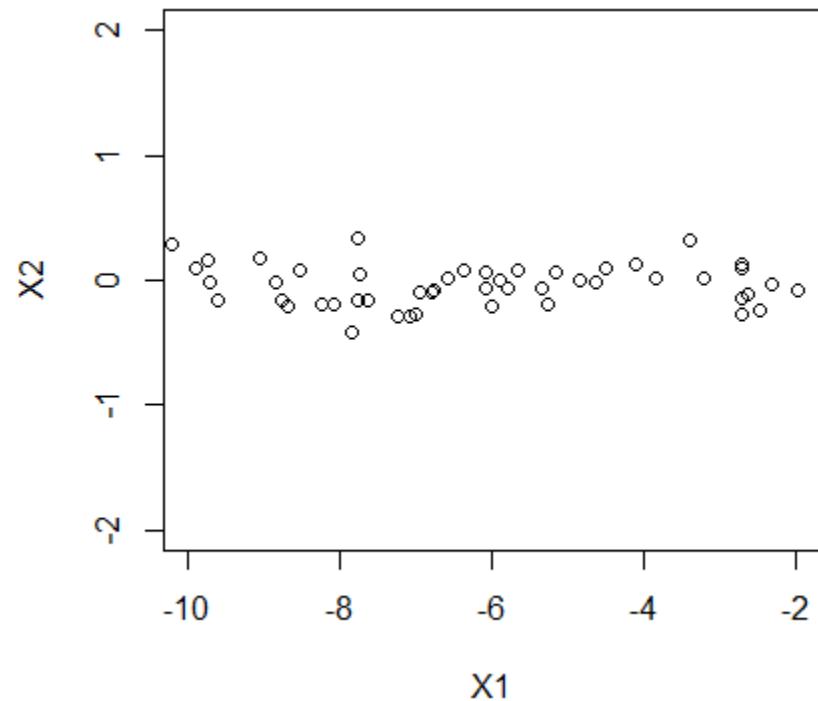
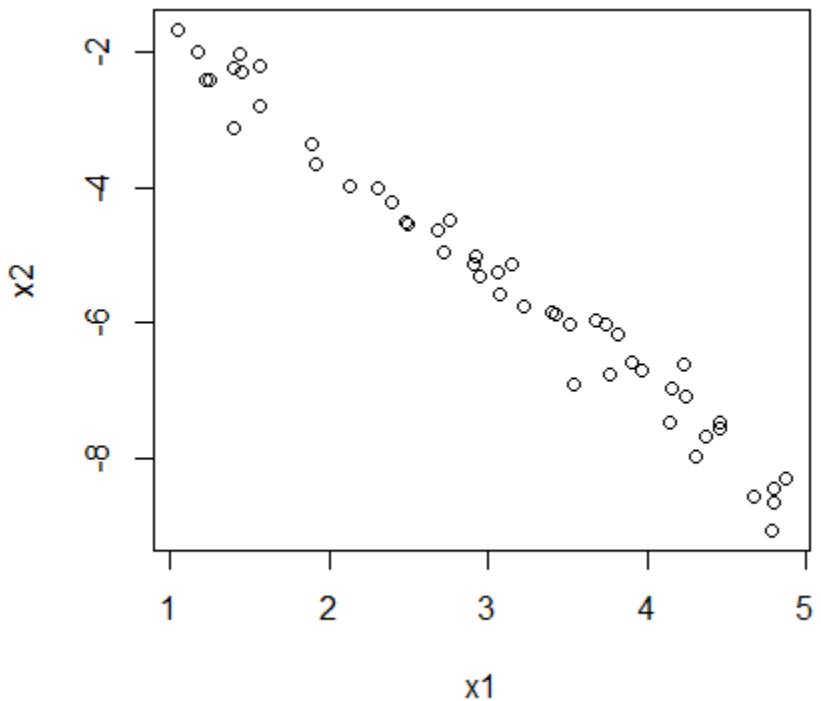
$\text{Var}(X1)=5.5134$

```
> eigen(c)
$values
[1] 5.51340440 0.02656068
```

$\text{Var}(X2)=0.02656$

Remember: The eigenvalues are exactly the same values!

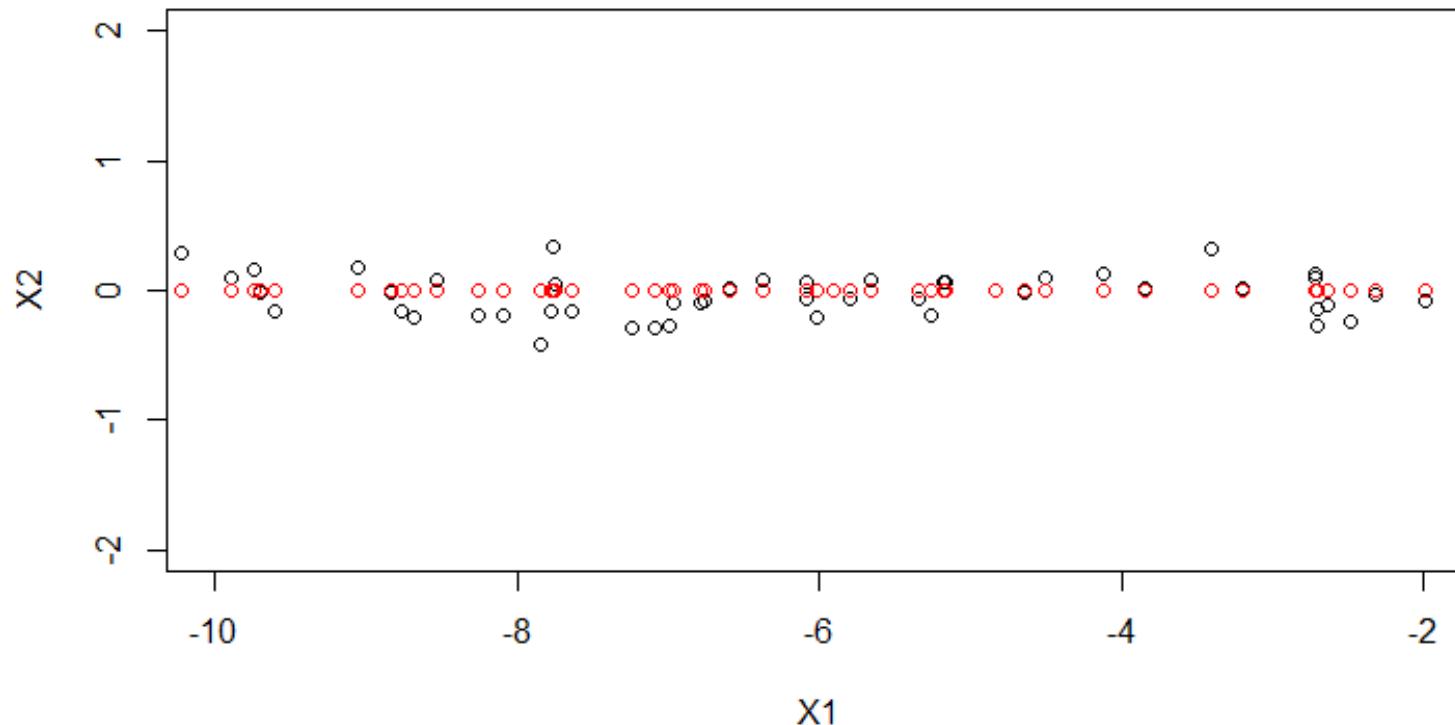
Data Transformed into the Basis of Principal Components



Since X_2 variable has small variance, we can now drop it by setting it to zero.

Dimensionality Reduction

X2 has been set to zero. Red dots are the approximated points.



So now instead of doing regression of y vs (x_1, x_2) , we are going to do regression of y vs X_1 . This is the point of doing PCA! It allows us to ignore variables with low variance.

**Lets now apply the idea to an example
with more features.**

CSE 7302c

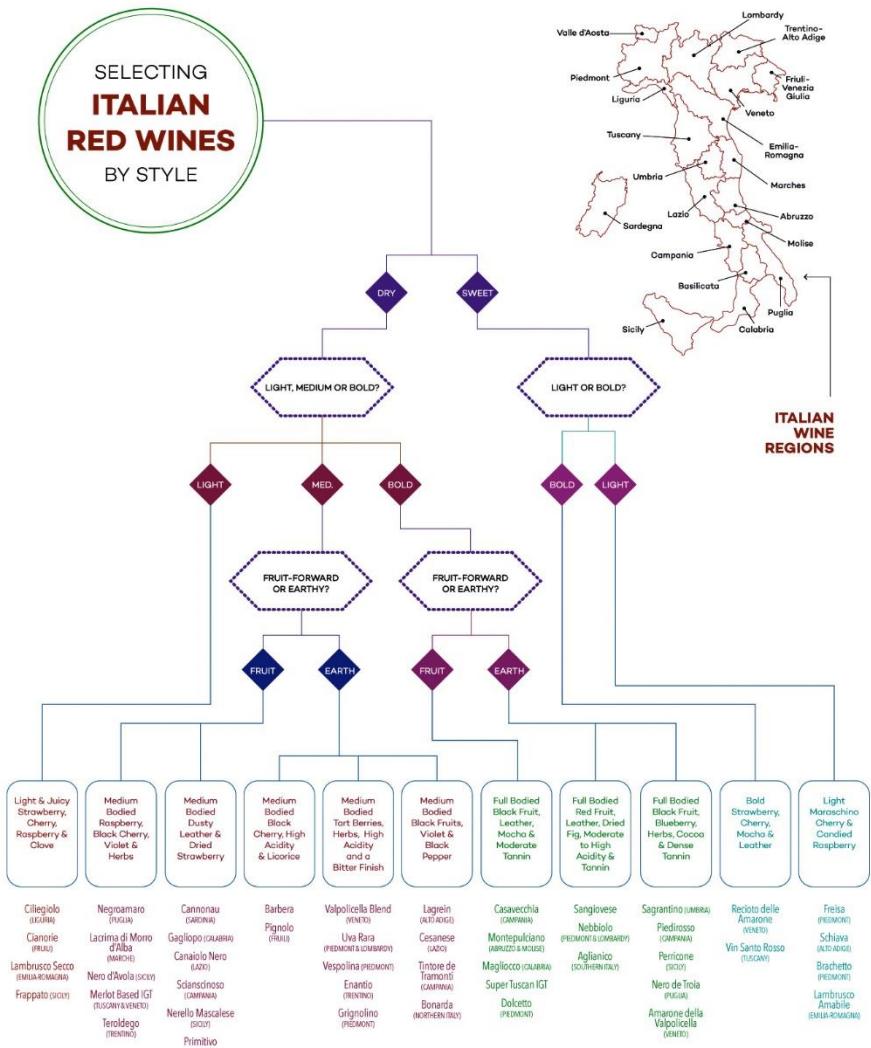




2C



Chemical Analysis of Wine



We have chemical analysis of 178 different wines produced by vineyards in Italy.

The wines come from 3 different grapes – Barolo, Grignolino, Barbera.

We are trying to identify the type of grape from 13 features.

Data Structure

WineClass	Alcohol	MalicAcid	Ash	AlkalinityAsh	Mg	Phenols	Flavanoids	NonFlavanoidPhenols	Proanthocyanins	Color	Hue	ODRatio	Proline
barolo	14.23	1.71	2.43	15.6	127	2.8	3.06	0.28	2.29	5.64	1.04	3.92	1065
barolo	13.2	1.78	2.14	11.2	100	2.65	2.76	0.26	1.28	4.38	1.05	3.4	1050
barolo	13.16	2.36	2.67	18.6	101	2.8	3.24	0.3	2.81	5.68	1.03	3.17	1185
barolo	14.37	1.95	2.5	16.8	113	3.85	3.49	0.24	2.18	7.8	0.86	3.45	1480
barolo	13.24	2.59	2.87	21	118	2.8	2.69	0.39	1.82	4.32	1.04	2.93	735
barolo	14.2	1.76	2.45	15.2	112	3.27	3.39	0.34	1.97	6.75	1.05	2.85	1450
barolo	14.39	1.87	2.45	14.6	96	2.5	2.52	0.3	1.98	5.25	1.02	3.58	1290
barolo	14.06	2.15	2.61	17.6	121	2.6	2.51	0.31	1.25	5.05	1.06	3.58	1295
barolo	14.83	1.64	2.17	14	97	2.8	2.98	0.29	1.98	5.2	1.08	2.85	1045
barolo	13.86	1.35	2.27	16	98	2.98	3.15	0.22	1.85	7.22	1.01	3.55	1045
barolo	14.1	2.16	2.3	18	105	2.95	3.32	0.22	2.38	5.75	1.25	3.17	1510
barolo	14.12	1.48	2.32	16.8	95	2.2	2.43	0.26	1.57	5	1.17	2.82	1280
barolo	13.75	1.73	2.41	16	89	2.6	2.76	0.29	1.81	5.6	1.15	2.9	1320
barolo	14.75	1.73	2.39	11.4	91	3.1	3.69	0.43	2.81	5.4	1.25	2.73	1150
barolo	14.38	1.87	2.38	12	102	3.3	3.64	0.29	2.96	7.5	1.2	3	1547
barolo	13.63	1.81	2.7	17.2	112	2.85	2.91	0.3	1.46	7.3	1.28	2.88	1310
barolo	14.3	1.92	2.72	20	120	2.8	3.14	0.33	1.97	6.2	1.07	2.65	1280
barolo	13.83	1.57	2.62	20	115	2.95	3.4	0.4	1.72	6.6	1.13	2.57	1130
barolo	14.19	1.59	2.48	16.5	108	3.3	3.93	0.32	1.86	8.7	1.23	2.82	1680
barolo	13.64	3.1	2.56	15.2	116	2.7	3.03	0.17	1.66	5.1	0.96	3.36	845
barolo	14.06	1.63	2.28	16	126	3	3.17	0.24	2.1	5.65	1.09	3.71	780
barolo	12.93	3.8	2.65	18.6	102	2.41	2.41	0.25	1.98	4.5	1.03	3.52	770
barolo	13.71	1.86	2.36	16.6	101	2.61	2.88	0.27	1.69	3.8	1.11	4	1035
barolo	12.85	1.6	2.52	17.8	95	2.48	2.37	0.26	1.46	3.93	1.09	3.63	1015
barolo	13.5	1.81	2.61	20	96	2.53	2.61	0.28	1.66	3.52	1.12	3.82	845
barolo	13.05	2.05	3.22	25	124	2.63	2.68	0.47	1.92	3.58	1.13	3.2	830

Dimension: 178 X 14.

Number of Features = 13

CSE 7302C



PCA Steps - R

- PCA is done only on predictor numerical variables
- Convert categories to dummy numerical variables
- If there are values of different order of magnitude, scale them.

```
> WineData <- read.csv("wine.csv")
> str(WineData)
'data.frame': 178 obs. of 14 variables:
 $ WineClass      : Factor w/ 3 levels "barbera","barolo",...: 2 2 2 2 2 2 2 2 2 ...
 $ Alcohol        : num 14.2 13.2 13.2 14.4 13.2 ...
 $ MalicAcid      : num 1.71 1.78 2.36 1.95 2.59 1.76 1.87 2.15 1.64 1.35 ...
 $ Ash            : num 2.43 2.14 2.67 2.5 2.87 2.45 2.45 2.61 2.17 2.27 ...
 $ AlkalinityAsh  : num 15.6 11.2 18.6 16.8 21 15.2 14.6 17.6 14 16 ...
 $ Mg             : int 127 100 101 113 118 112 96 121 97 98 ...
 $ Phenols        : num 2.8 2.65 2.8 3.85 2.8 3.27 2.5 2.6 2.8 2.98 ...
 $ Flavanoids     : num 3.06 2.76 3.24 3.49 2.69 3.39 2.52 2.51 2.98 3.15 ...
 $ NonFlavanoidPhenols: num 0.28 0.26 0.3 0.24 0.39 0.34 0.3 0.31 0.29 0.22 ...
 $ Proanthocyanins: num 2.29 1.28 2.81 2.18 1.82 1.97 1.98 1.25 1.98 1.85 ...
 $ Color          : num 5.64 4.38 5.68 7.8 4.32 6.75 5.25 5.05 5.2 7.22 ...
 $ Hue            : num 1.04 1.05 1.03 0.86 1.04 1.05 1.02 1.06 1.08 1.01 ...
 $ ODRatio        : num 3.92 3.4 3.17 3.45 2.93 2.85 3.58 3.58 2.85 3.55 ...
 $ Proline        : int 1065 1050 1185 1480 735 1450 1290 1295 1045 1045 ...'

> # PCA analysis is done only on the predictors
> wine.predictors <- WineData[,-1]
>
> # Since the predictors are of completely different magnitude,
> # we need scale them before the analysis.
> scaled.Predictors <- scale(wine.predictors)
'
```

Principal Components in R

- *princomp* performs a principal components analysis on a given numeric data matrix
- Its outputs are:
 - *loadings*: The coefficients for linear transformations into the new variables
 - *scores*: The input features into the transformed bases
 - It also gives a summary of percentage of variance explained by each Principal Component

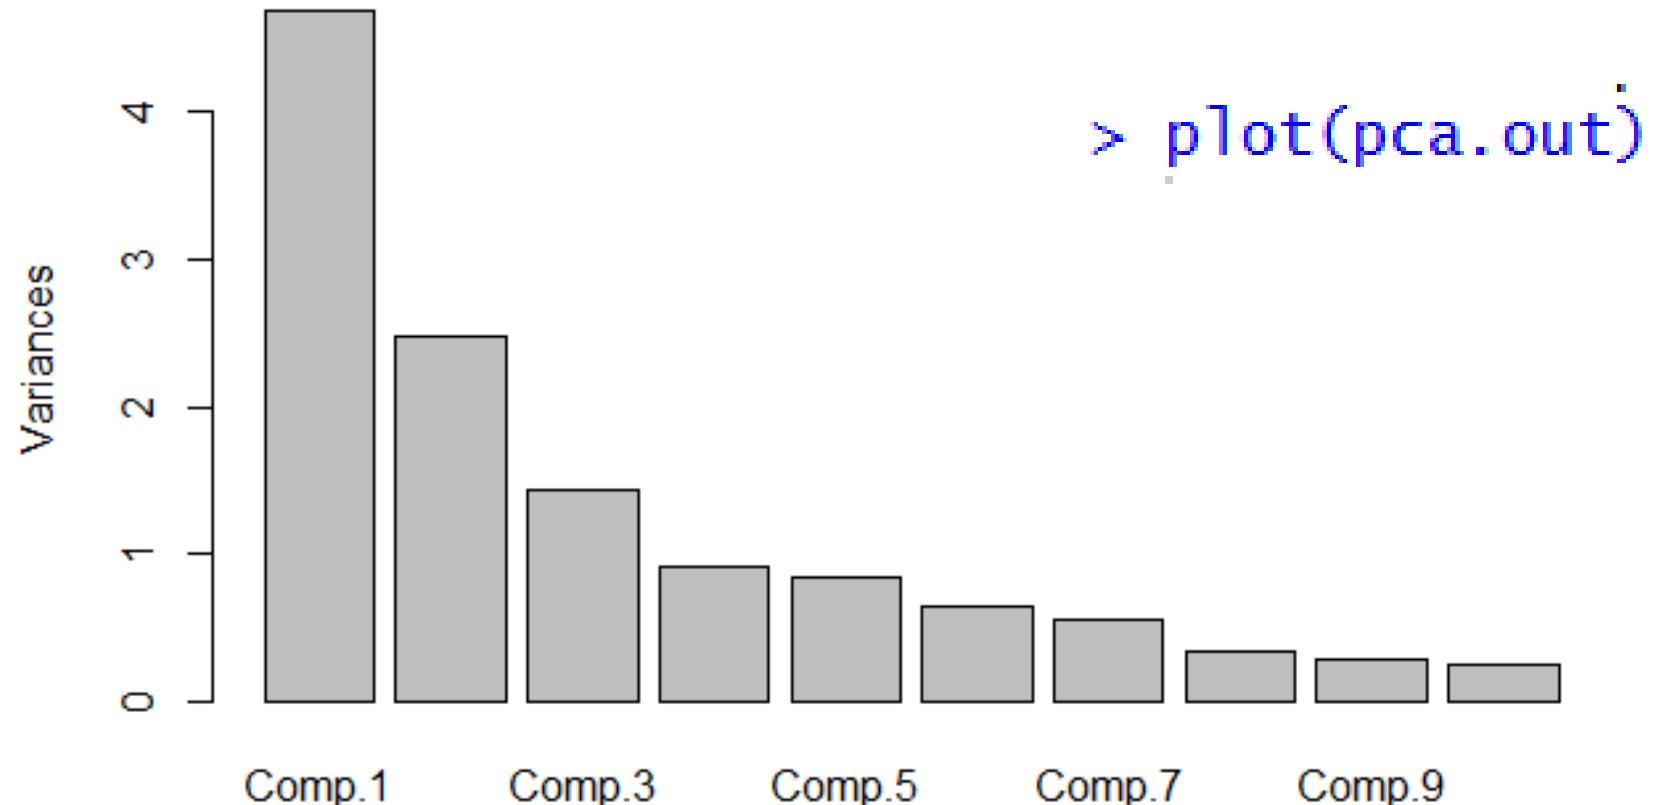
```

> # compute PCs
> pca.out = princomp(scaled.Predictors)
> summary(pca.out)
Importance of components:
                                         Comp.1    Comp.2    Comp.3    Comp.4    Comp.5    Comp.6    Comp.7
Standard deviation     2.1631951 1.5757366 1.1991447 0.9559347 0.92110518 0.79878171 0.74022473
Proportion of Variance 0.3619885 0.1920749 0.1112363 0.0706903 0.06563294 0.04935823 0.04238679
Cumulative Proportion  0.3619885 0.5540634 0.6652997 0.7359900 0.80162293 0.85098116 0.89336795
                                         Comp.8    Comp.9    Comp.10   Comp.11   Comp.12   Comp.13
Standard deviation     0.58867607 0.53596364 0.49949266 0.47383559 0.40966094 0.320619963
Proportion of Variance 0.02680749 0.02222153 0.01930019 0.01736836 0.01298233 0.007952149
Cumulative Proportion  0.92017544 0.94239698 0.96169717 0.97906553 0.99204785 1.0000000000

```

- There are 13 principal components (since there were 13 initial features).
- The output principal components are ordered according to the percentage of variance explained.
- The top row gives the (sqrt of) eigenvalues, the 2nd row gives the percentage of explained variance.
- First 8 components explain over 90% of the variance!

pca.out



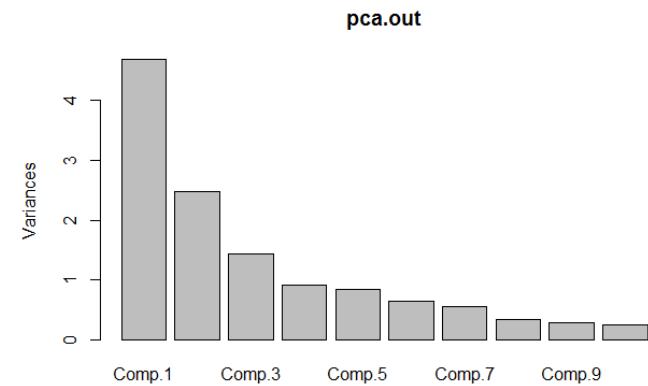
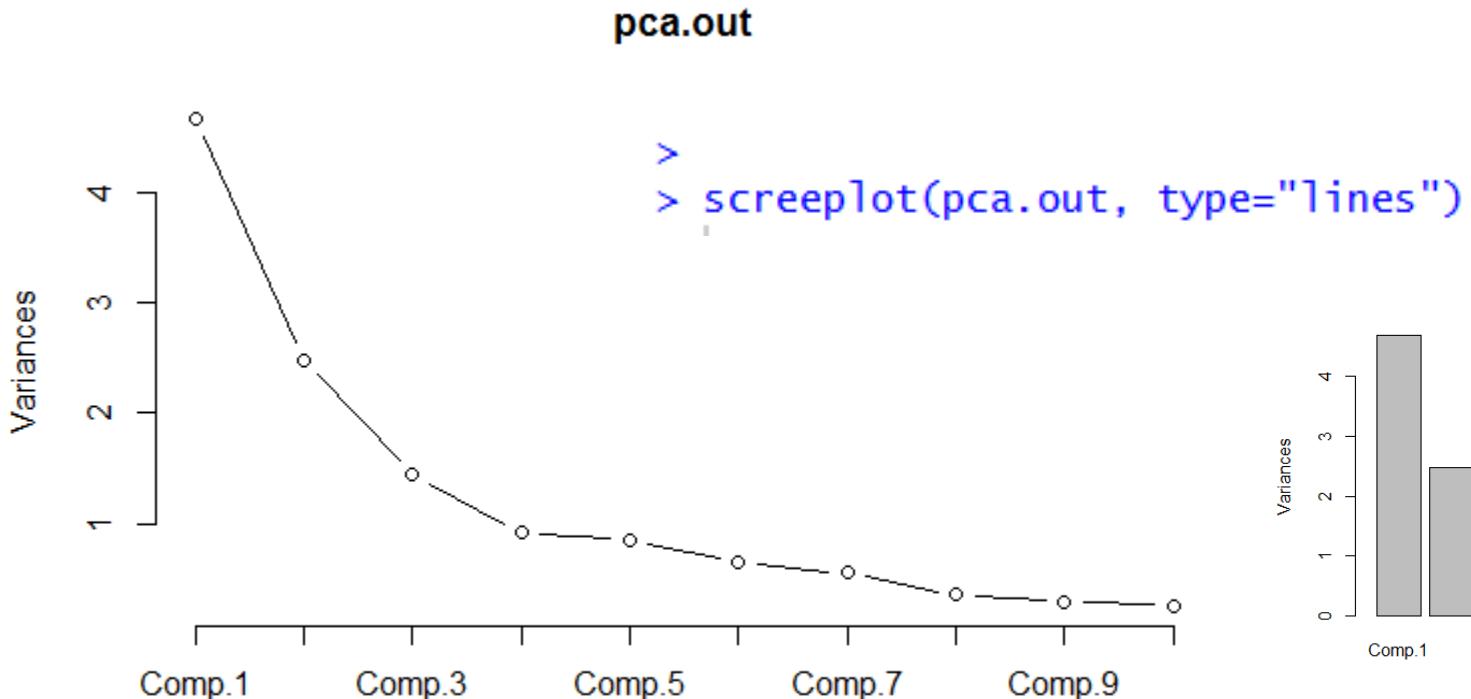
```
> plot(pca.out)
```

Dimensionality Reduction

- Let us say we set a threshold of 80% variance explanatory power.
- Keeping the first 5 principal components explains 80% of variance.
- So, we can choose to ignore the remaining 8 components and build our model.

```
>  
> compressed_features = pca$scores[,1:5]  
> library(nnet)  
> multout.pca <- multinom(wineData$wine.class ~ compressed_features)
```

How Many Components to Choose?



Rule of thumb: Pick the point, which is at the “elbow” in the plot.
Over here, it happens at around 4th or 5th component.

CSE 7302C



PCA Performance: Multinomial Classification

Regular With PCA

```
> summary(multout.full)
Call:
multinom(formula = WineClass ~ scaled.Predictors, data = WineData)
```

Coefficients:

	(Intercept)	scaled.PredictorsAlcohol	scaled.PredictorsMalicAcid	scaled.PredictorsAsh	scaled.PredictorsAlkalinityAsh
barolo	-0.5736	11.552	-1.888	5.39	-13.075
grignolino	4.8870	-7.864	-3.600	-11.81	5.087
	scaled.PredictorsMg	scaled.PredictorsPhenols	scaled.PredictorsFlavanoids	scaled.PredictorsNonFlavanoidPhenols	
barolo	2.5506	8.113	11.892		-5.154
grignolino	0.3456	-1.243	9.796		4.570
	scaled.PredictorsProanthocyanins	scaled.PredictorsColor	scaled.PredictorsHue	scaled.PredictorsODRatio	
barolo		3.339	0.9285	5.60	9.196
grignolino		2.597	-19.8481	12.53	5.140
	scaled.PredictorsProline				
barolo		14.49			
grignolino		-11.35			

Std. Errors:

	(Intercept)	scaled.PredictorsAlcohol	scaled.PredictorsMalicAcid	scaled.PredictorsAsh	scaled.PredictorsAlkalinityAsh
barolo	119427	167686	58547	109387	106846
grignolino	12729	4156	1606	8713	3680
	scaled.PredictorsMg	scaled.PredictorsPhenols	scaled.PredictorsFlavanoids	scaled.PredictorsNonFlavanoidPhenols	
barolo	79264	151221	80842		137785
grignolino	7269	4339	8241		3672
	scaled.PredictorsProanthocyanins	scaled.PredictorsColor	scaled.PredictorsHue	scaled.PredictorsODRatio	
barolo		82716	114552	74901	135373
grignolino		5388	7194	5822	3090
	scaled.PredictorsProline				
barolo		99521			
grignolino		10364			

Residual Deviance: 0.0001846
AIC: 56

AIC = 56

```
> summary(multout.pca)
Call:
multinom(formula = WineData$WineClass ~ compressed_features)
```

Coefficients:

	(Intercept)	compressed_featuresComp.1	compressed_featuresComp.2	compressed_featuresComp.3	compressed_featuresComp.4	
barolo	11.21		-33.58	3.355	3.839	1.303
grignolino	19.30		-14.06	44.527	14.878	2.494
	compressed_featuresComp.5					
barolo		20.489				
grignolino		7.312				
	Std. Errors:					
	(Intercept)	compressed_featuresComp.1	compressed_featuresComp.2	compressed_featuresComp.3	compressed_featuresComp.4	
barolo	647.5		395.65	262.6	44.72	89.35
grignolino	258.6		36.12	394.5	165.18	78.81
	compressed_featuresComp.5					
barolo		294.0				
grignolino		127.4				
	Residual Deviance:	0.07747				
	AIC:	24.08				

AIC = 24.08

AIC in Multinomial Logistic Regression is given by $AIC = D + 2kr$, where k is the # of parameters (including intercept) and r is the # of categories minus 1.

PCA does not always guarantee better performance. The main use is dimensionality reduction.

CSE 7302c

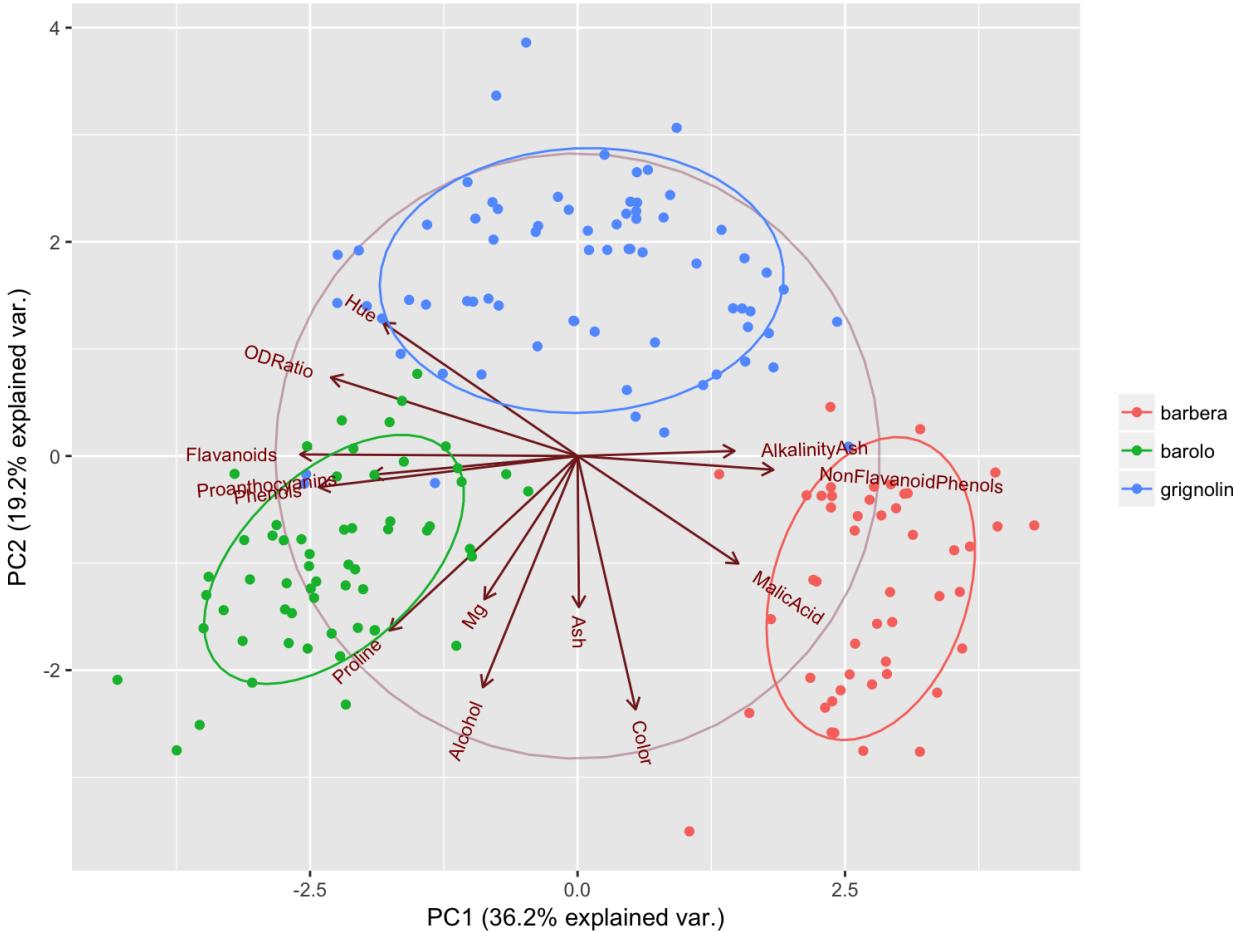


PCA: Visualization

- Many times visualizing the spread in the data-points using only the first 2 Principal Components can provide useful information about the spread of the data
- The function *biplot* lets you do that
- Also check *ggbiplot* in the *library(ggbiplot)*

For installing ggbiplot see: <http://www.vince.vu/software/#ggbiplot>

Projection Using First 2 Components



Each class nicely gets segmented when viewed from the Principal Component basis.

biplot is also occasionally used to identify outliers.

```
> library(ggbiplot)
> g <- ggbiplot(pca.out, obs.scale = 1, var.scale = 1,
+                 groups = WineData$WineClass, ellipse = TRUE
, circle = TRUE)
> g + scale_color_discrete(name = '')
```

1000 Genomes Single Chromosome PCA Example

- Computing principal components of genomic variant data across one chromosome from 2504 people from the 1000 genomes project.
- Effective clustering of people by ethnicity (5 ethnicities) using PCA.
- A sparse matrix with 2504 rows (people) by 18,12,841 columns (variants).

PEL	Peruvians from Lima, Peru	AMR	1	1	1
GIH	Gujarati Indian from Houston, Texas	SAS	1	1	1
PJL	Punjabi from Lahore, Pakistan	SAS	1	1	1
BEB	Bengali from Bangladesh	SAS	1	1	1
STU	Sri Lankan Tamil from the UK	SAS	1	1	1
ITU	Indian Telugu from the UK	SAS	1	1	1

These populations have been divided into 5 super populations

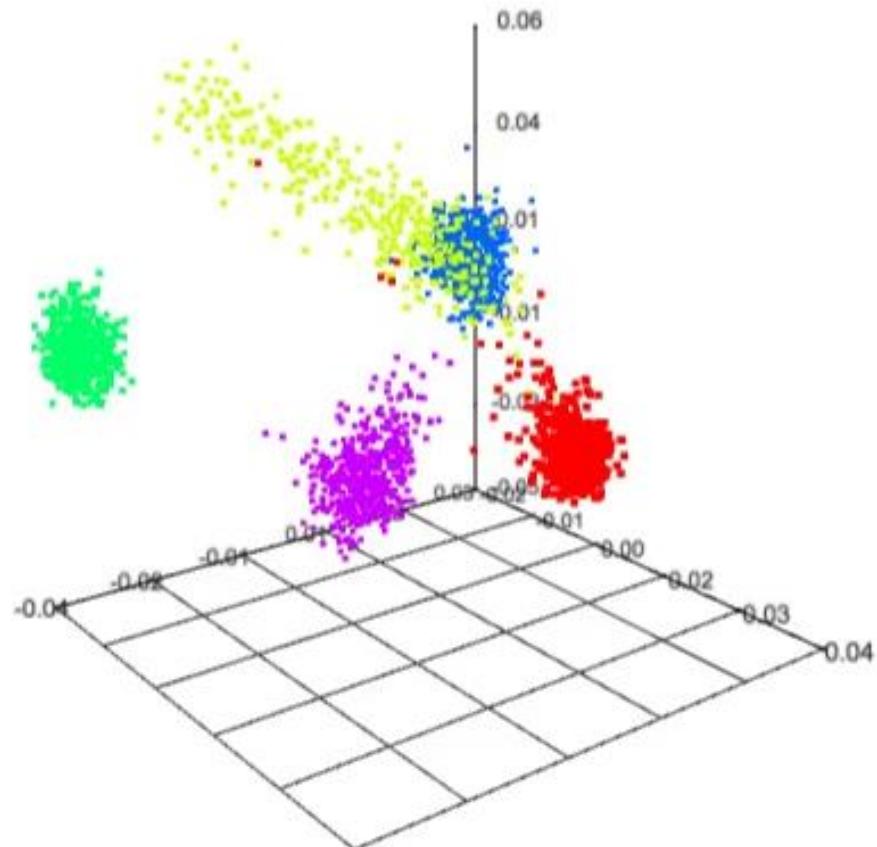
- **AFR**, African
- **AMR**, Ad Mixed American
- **EAS**, East Asian
- **EUR**, European
- **SAS**, South Asian

When the code **ALL** is used this means that all individuals from that release are being considered.

Source: https://bwlewis.github.io/1000_genomes_examples/PCA.html and
<http://www.internationalgenome.org/faq/which-populations-are-part-your-study/>

Last accessed: September 09, 2017

1000 Genomes Single Chromosome PCA Example



Source: https://bwlewis.github.io/1000_genomes_examples/PCA.html
Last accessed: September 09, 2017

CSE 7302c

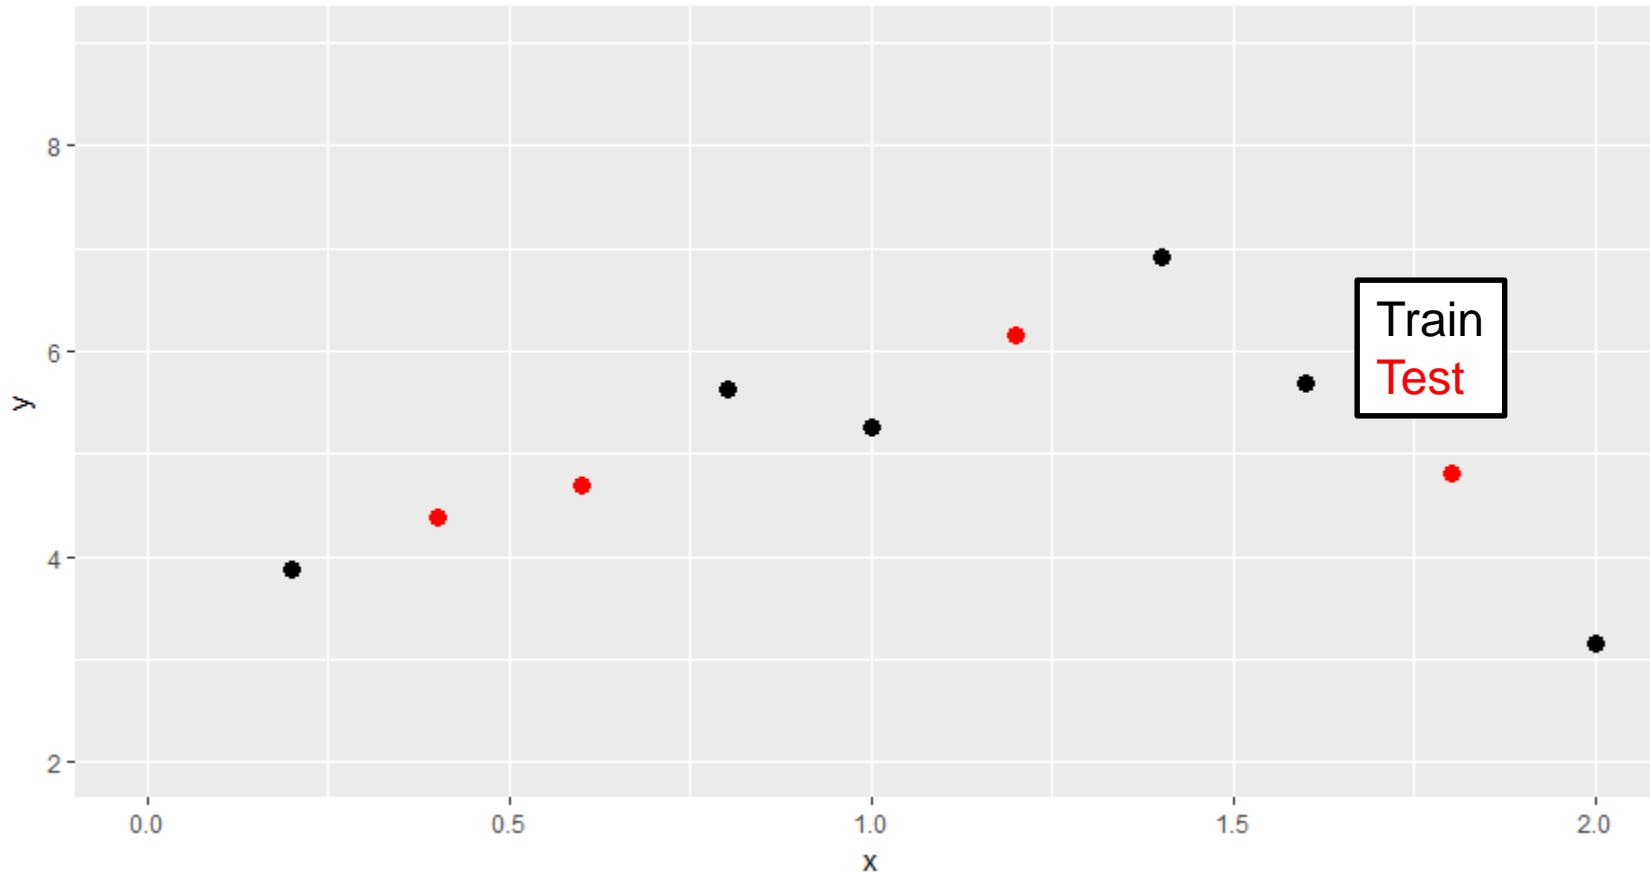


**Advanced Regression Methods: Ridge
Regression, LASSO Regression, Elastic Nets**

REGULARIZATION

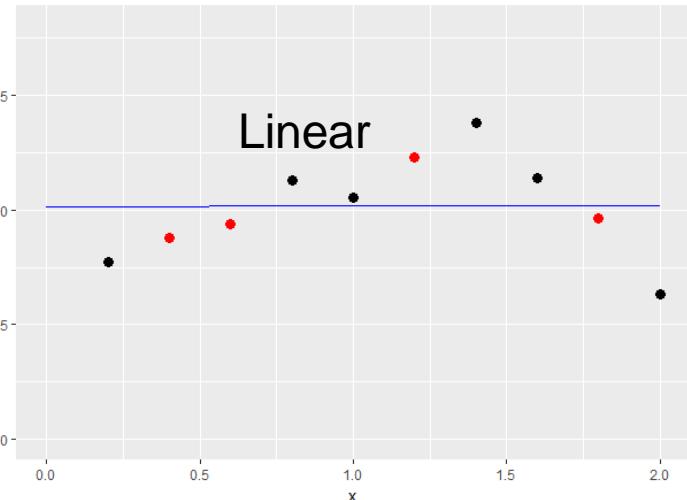
CSE 7302c





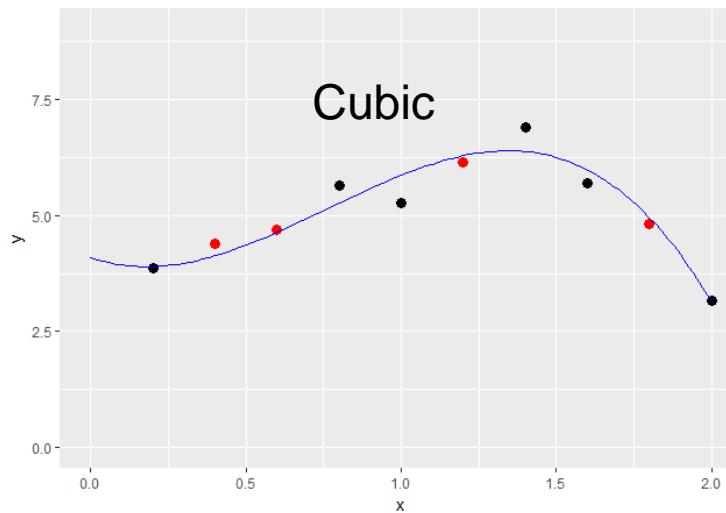
Decision point: What order polynomial should we use for the fit?

Problem with Overfitting

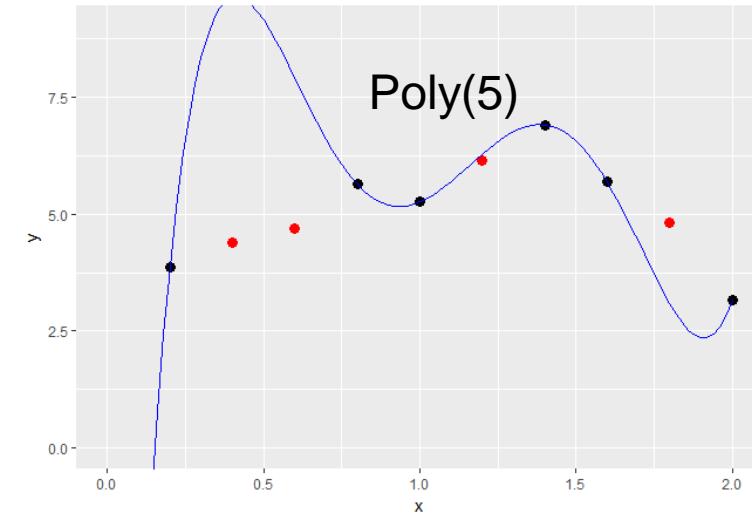


TrainErr=1.23
TestErr=0.7

TrainErr=0.38
TestErr=0.17



TrainErr=0
TestErr=3.2



Overfitting Problem

$$y = \beta_0 + \beta_1 x + \beta_2 x^2 + \beta_3 x^3 + \beta_4 x^4 + \dots$$

How many β_i do we need for a reasonable fit, while at the same time minimizing the risk of overfitting?

In general, given a regression problem with several features, is there a way to determine how many β_i do we need for a reasonable fit, while at the same time minimizing the risk of overfitting?

$$y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \beta_3 x_3 + \beta_4 x_4 + \dots$$

Regularization

- Standard Least Squares Regression involves minimizing SSE

$$\text{SSE} = \sum_{j=1}^N (y_j - \sum_{i=0}^d \beta_i \cdot x_i)^2 = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^T \cdot (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$$

\mathbf{y} = vector of all training responses y_j

\mathbf{X} = matrix of all training samples \mathbf{x}_j

- We can reduce overfitting by penalizing large coefficients (**must standardize** in order to be able to compare/add β s)

Ridge Regression

$$\min_{\beta} \left(\sum_{j=1}^N (y_j - \sum_{i=0}^d \beta_i \cdot x_i)^2 + \lambda \sum_{i=1}^d \beta_i^2 \right)$$

- Penalization of large coefficients through L2 norm (another name for Least Squares Error)
- Regularization parameter λ controls the extent of penalization
- $\lambda = 0$ corresponds to Least Squares Regression

Ridge Regression

- Occam's Razor

$$\min_{\beta} \left(\sum_{j=1}^N (y_j - \sum_{i=0}^d \beta_i \cdot x_i)^2 + \lambda \sum_{i=1}^d \beta_i^2 \right)$$

- Least Squares Regression

- Fails when the number of features p is larger than number of data points n , e.g., in genetics studies. For example, there are infinite solutions for $5 = \beta_1 * 3 + \beta_2 * 2$ (two unknowns and one equation)

CSE 7302c



Ridge Regression

- Least Squares Regression
 - Fails when we have perfect multicollinearity

$y = x_1 + x_2$ or $y = 0.5x_1 + 1.5x_2$ or $y = 2x_1$ and so on

$$\min_{\beta} \left(\sum_{j=1}^N (y_j - \sum_{i=0}^d \beta_i \cdot x_i)^2 + \lambda \sum_{i=1}^d \beta_i^2 \right)$$

y	x1	x2
0	0	0
2	1	1
4	2	2
6	3	3

- Ridge Regression allows for a unique solution even in above situation. Which equation will you choose if you add the constraint to minimize $\sum_{i=1}^d b_i^2$ as well?

Ridge Regularization

The penalty term “regularizes” the problem when there is perfect multicollinearity

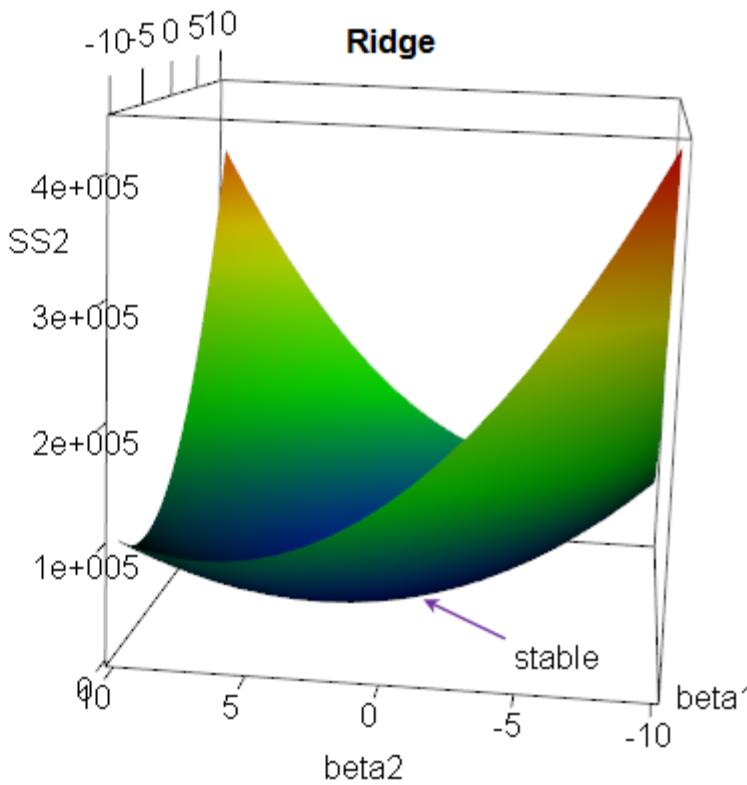
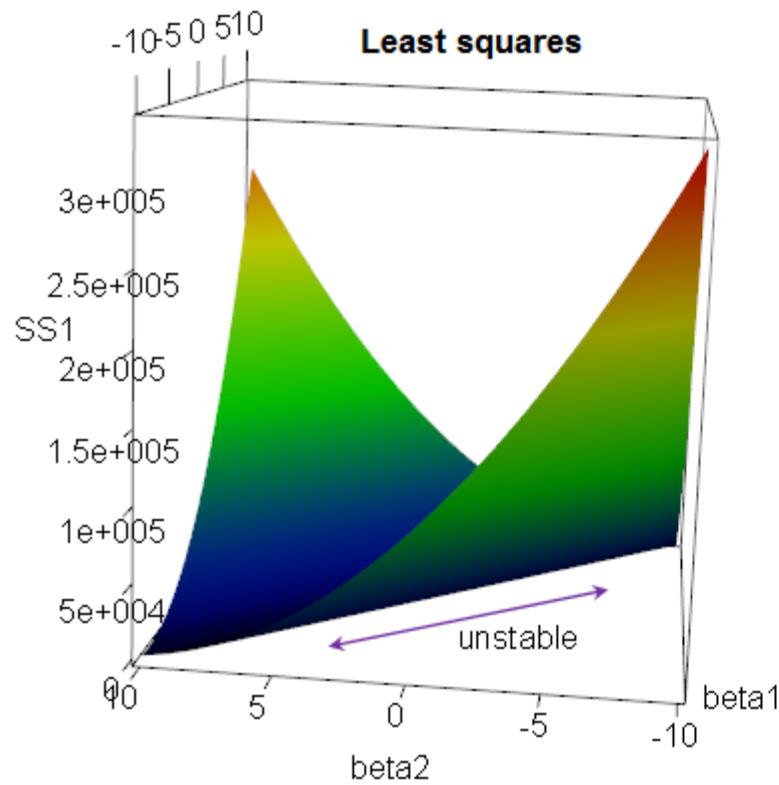
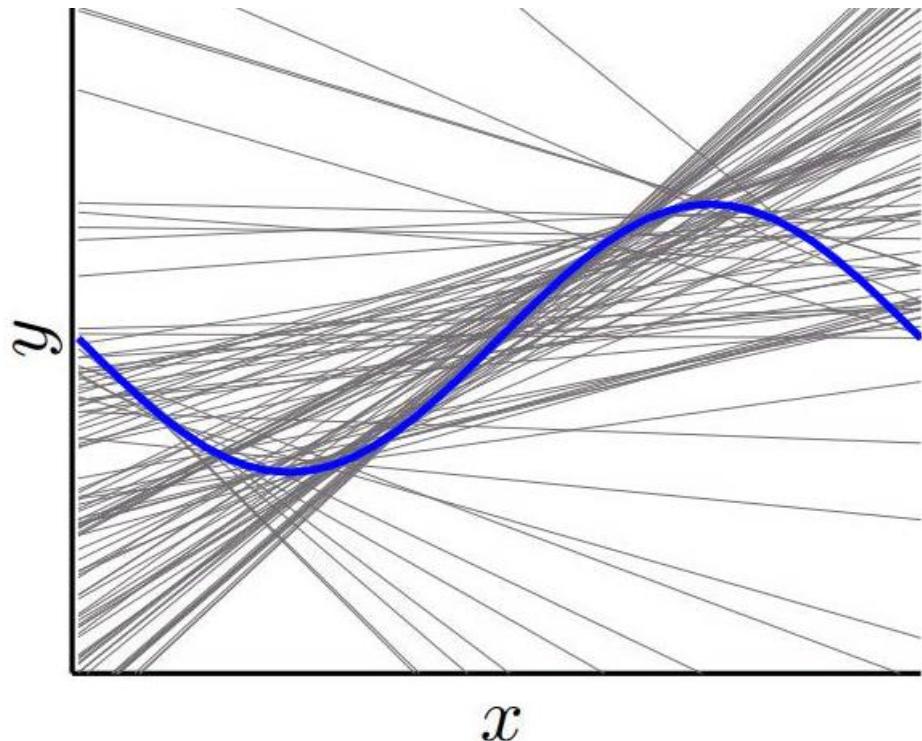
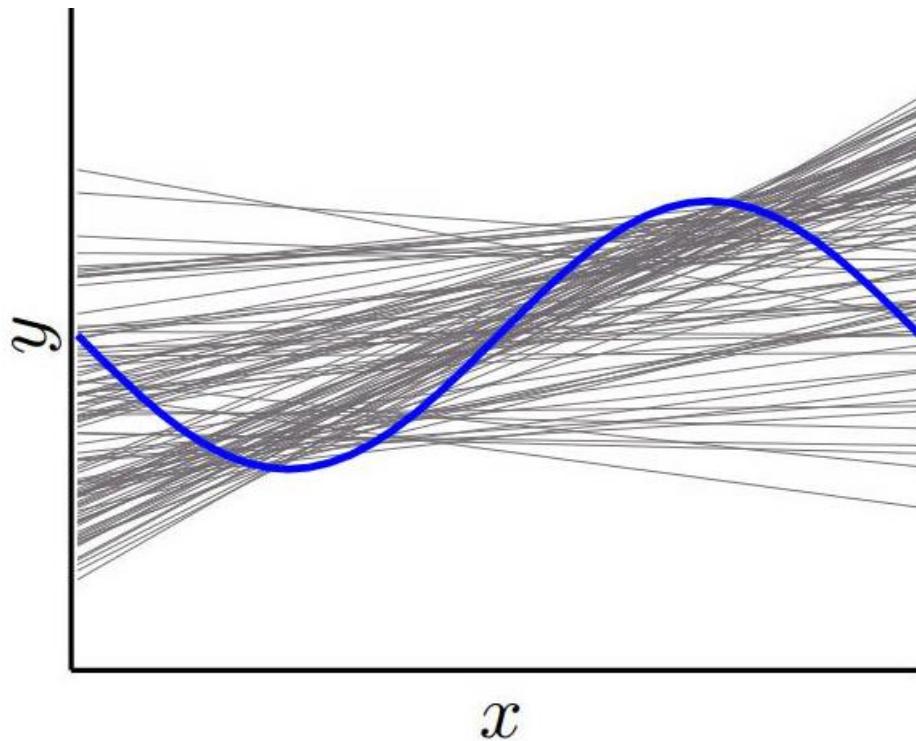


Image Source: <https://stats.stackexchange.com/questions/151304/why-is-ridge-regression-called-ridge-why-is-it-needed-and-what-happens-when>

Effect of Regularization



without regularization



with regularization

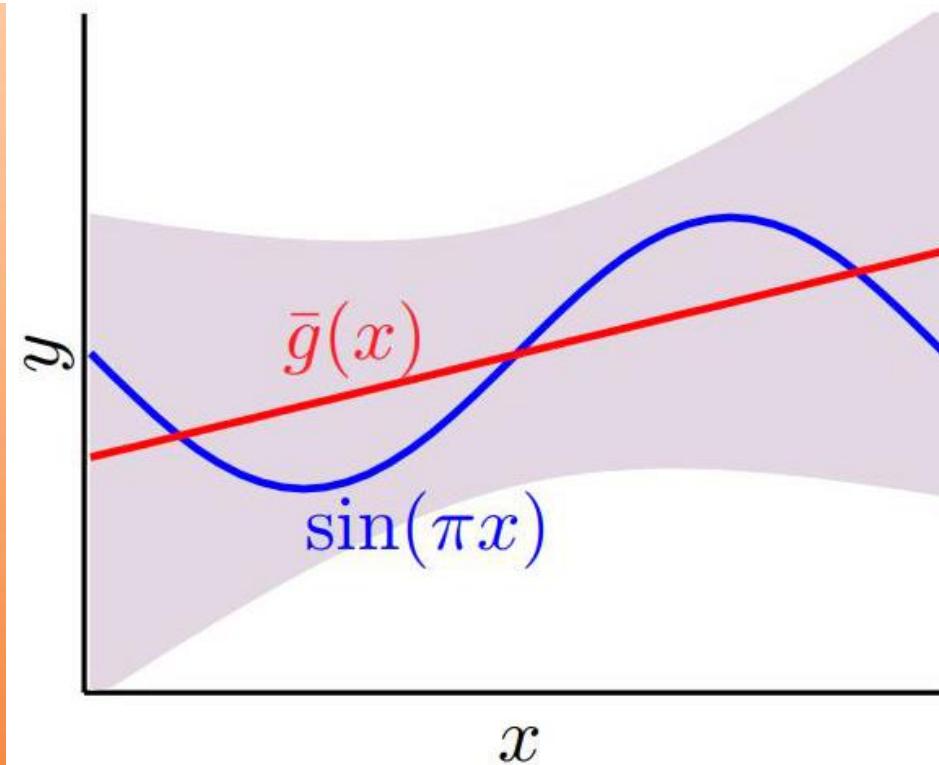
Source: <http://work.caltech.edu/slides/slides12.pdf>

CSE 7302c



Effect of Regularization

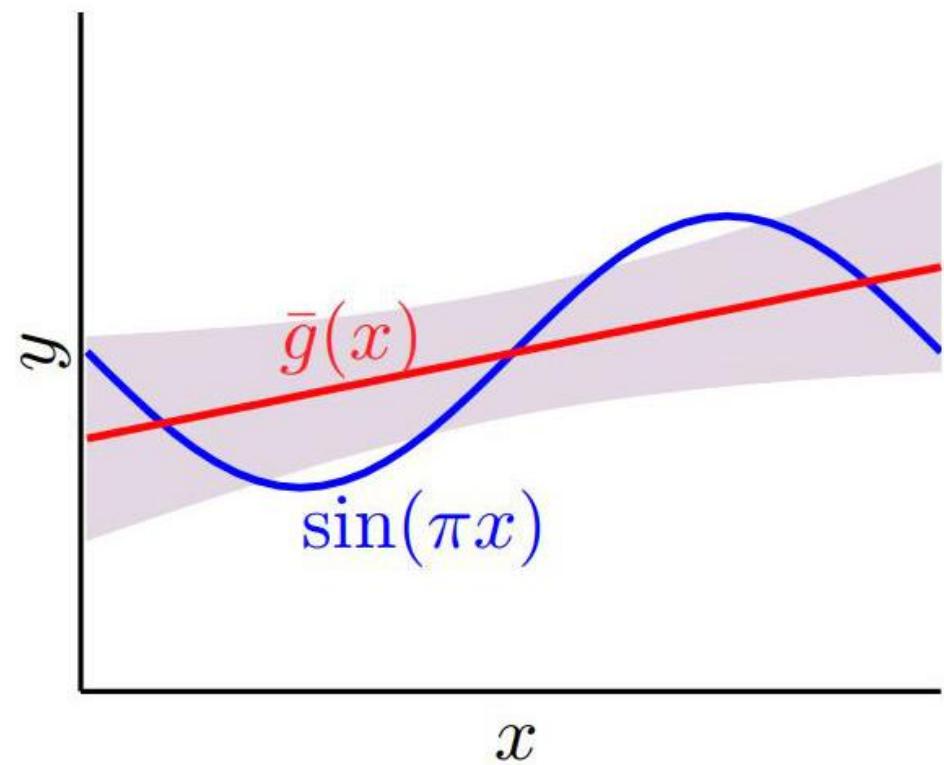
without regularization



bias = **0.21**

var = **1.69**

with regularization



bias = **0.23**

var = **0.33**

Source: <http://work.caltech.edu/slides/slides12.pdf>

CSE 7302c
INTERNATIONAL SCHOOL OF ENGINEERING
INSOFE



LASSO Regularization

$$\min_{\beta} \left(\sum_{j=1}^N (y_j - \sum_{i=0}^d \beta_i \cdot x_i)^2 + \lambda \sum_{i=1}^d |\beta_i| \right)$$

- Penalization of large coefficients through L1 norm (another name for Least Absolute Error)
- Least Absolute Shrinkage and Selection Operator (LASSO)

CSE 7302c



LASSO Regularization

$$\min_{\beta} \left(\sum_{j=1}^N (y_j - \sum_{i=0}^d \beta_i \cdot x_i)^2 + \lambda \sum_{i=1}^d |\beta_i| \right)$$

y	x1	x2
0	0	0
2	1	1
4	2	2
6	3	3

Not useful when there is strong multicollinearity.

$y = x_1 + x_2$ or $y = 0.5x_1 + 1.5x_2$ or $y = 2x_1$ and so on

CSE 7302c



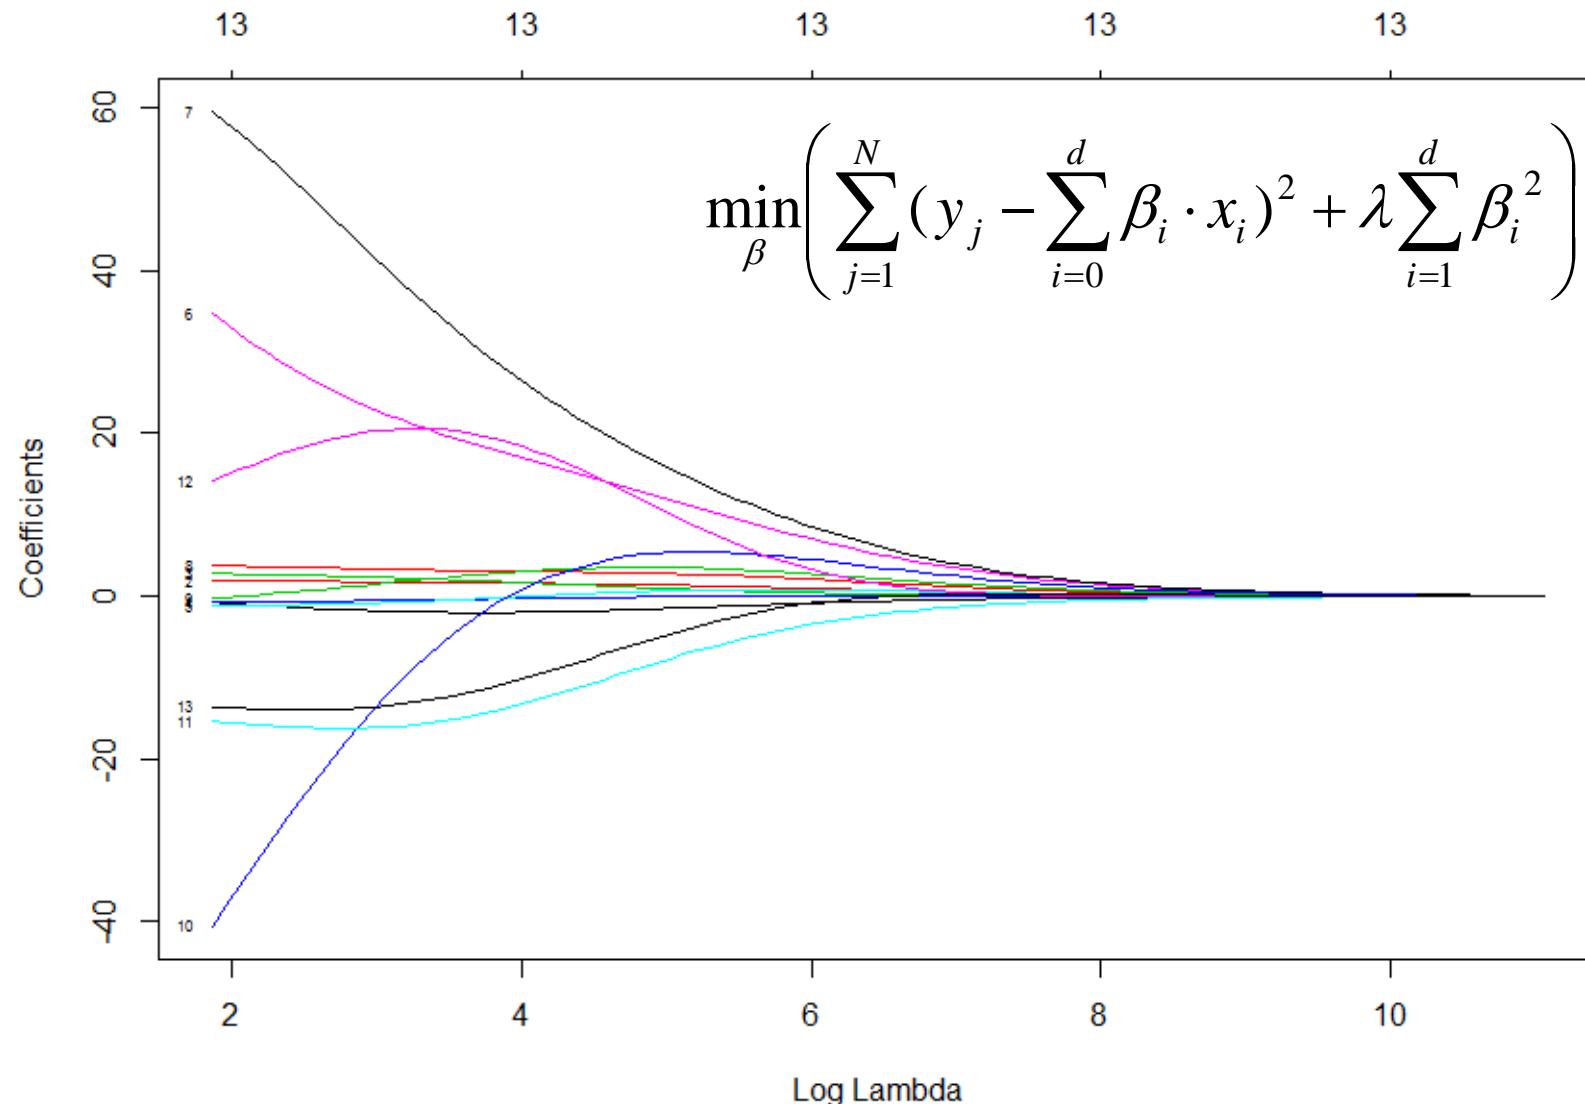
Regularization – Toy Purchase Data



CSE 7302c



Ridge: Coefficients



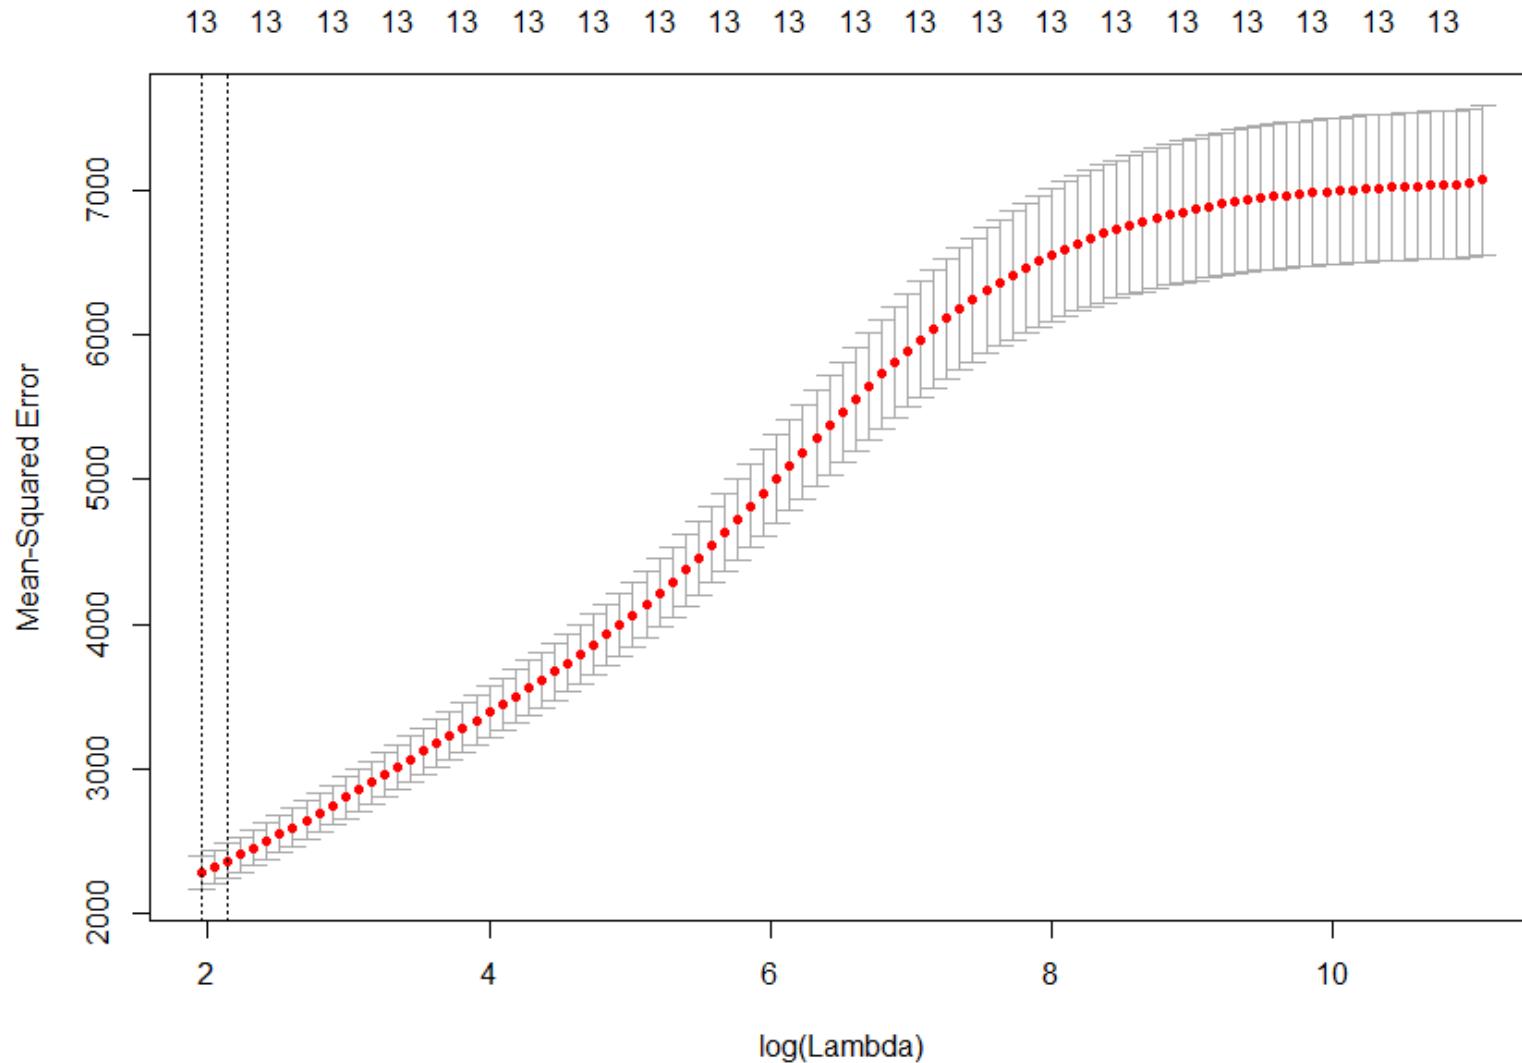
How to choose λ

- Choose a value for λ .
- Run k -fold cross validation and calculate mean “squared” error.
- Repeat the above 2 steps for various λ values.
- Pick λ that gives least error. R recommends the range from $\min \lambda$ to 1 stdev away.
- Pick λ by looking at the coefficients corresponding to these λ values, in addition to the minimum error.

CSE 7302c

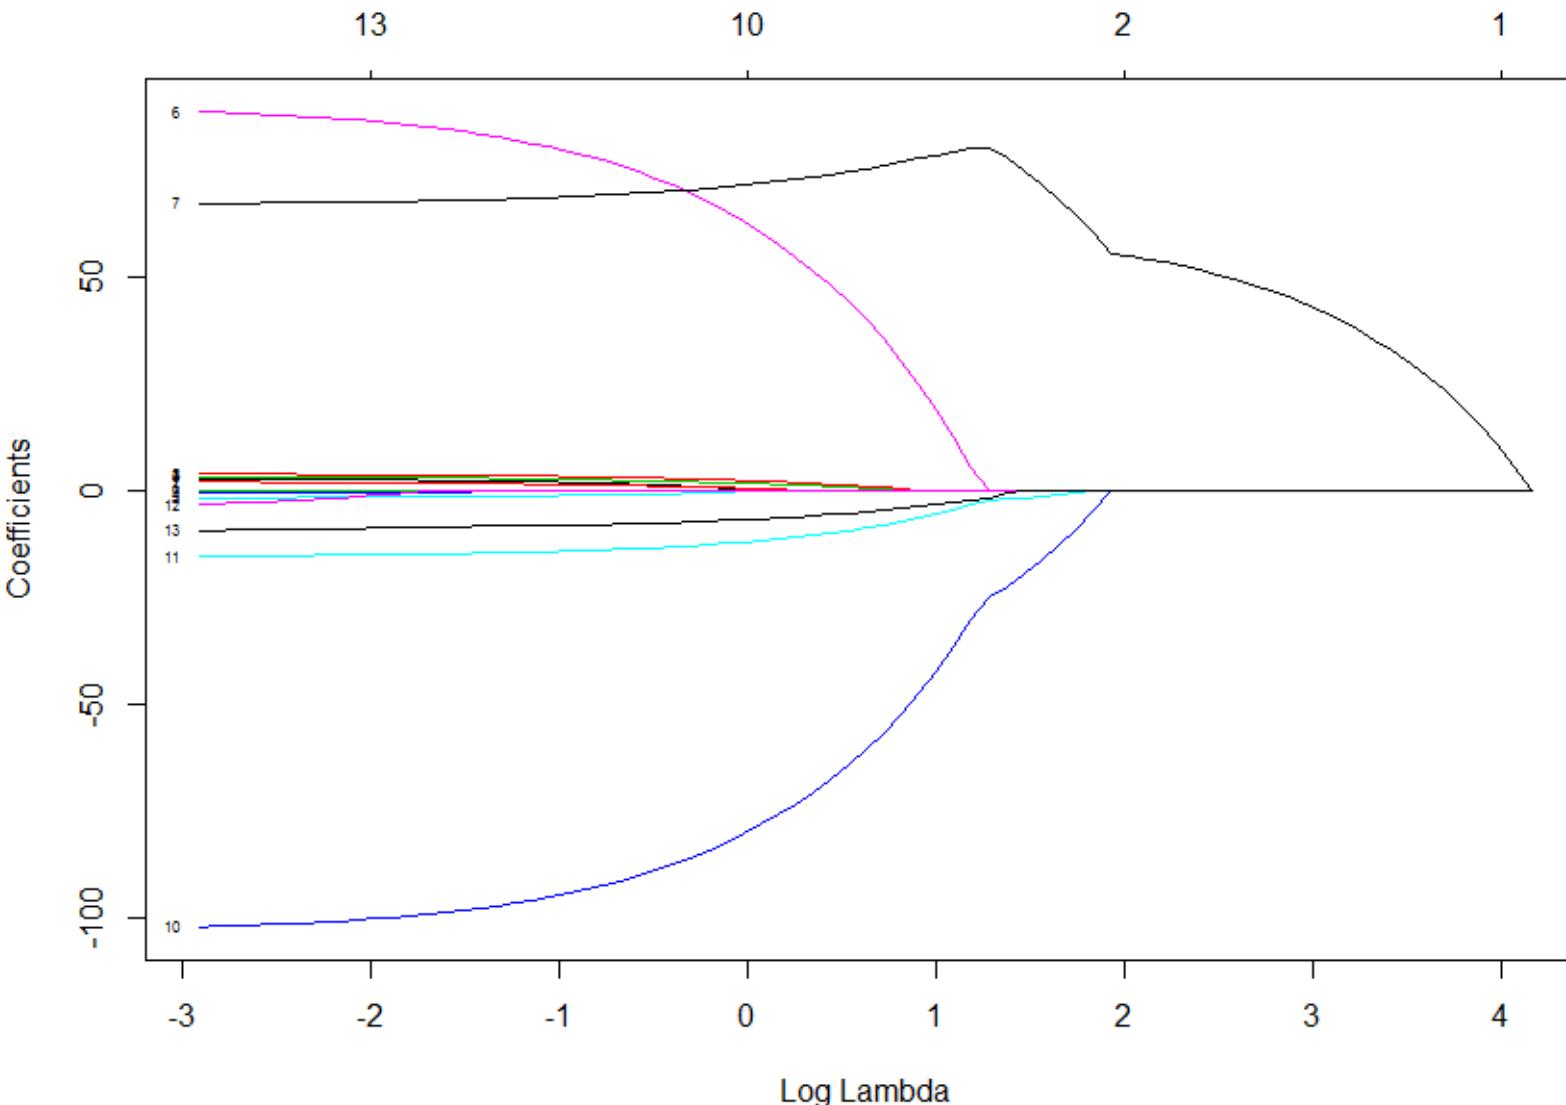


Ridge: MSE



LASSO: Coefficients

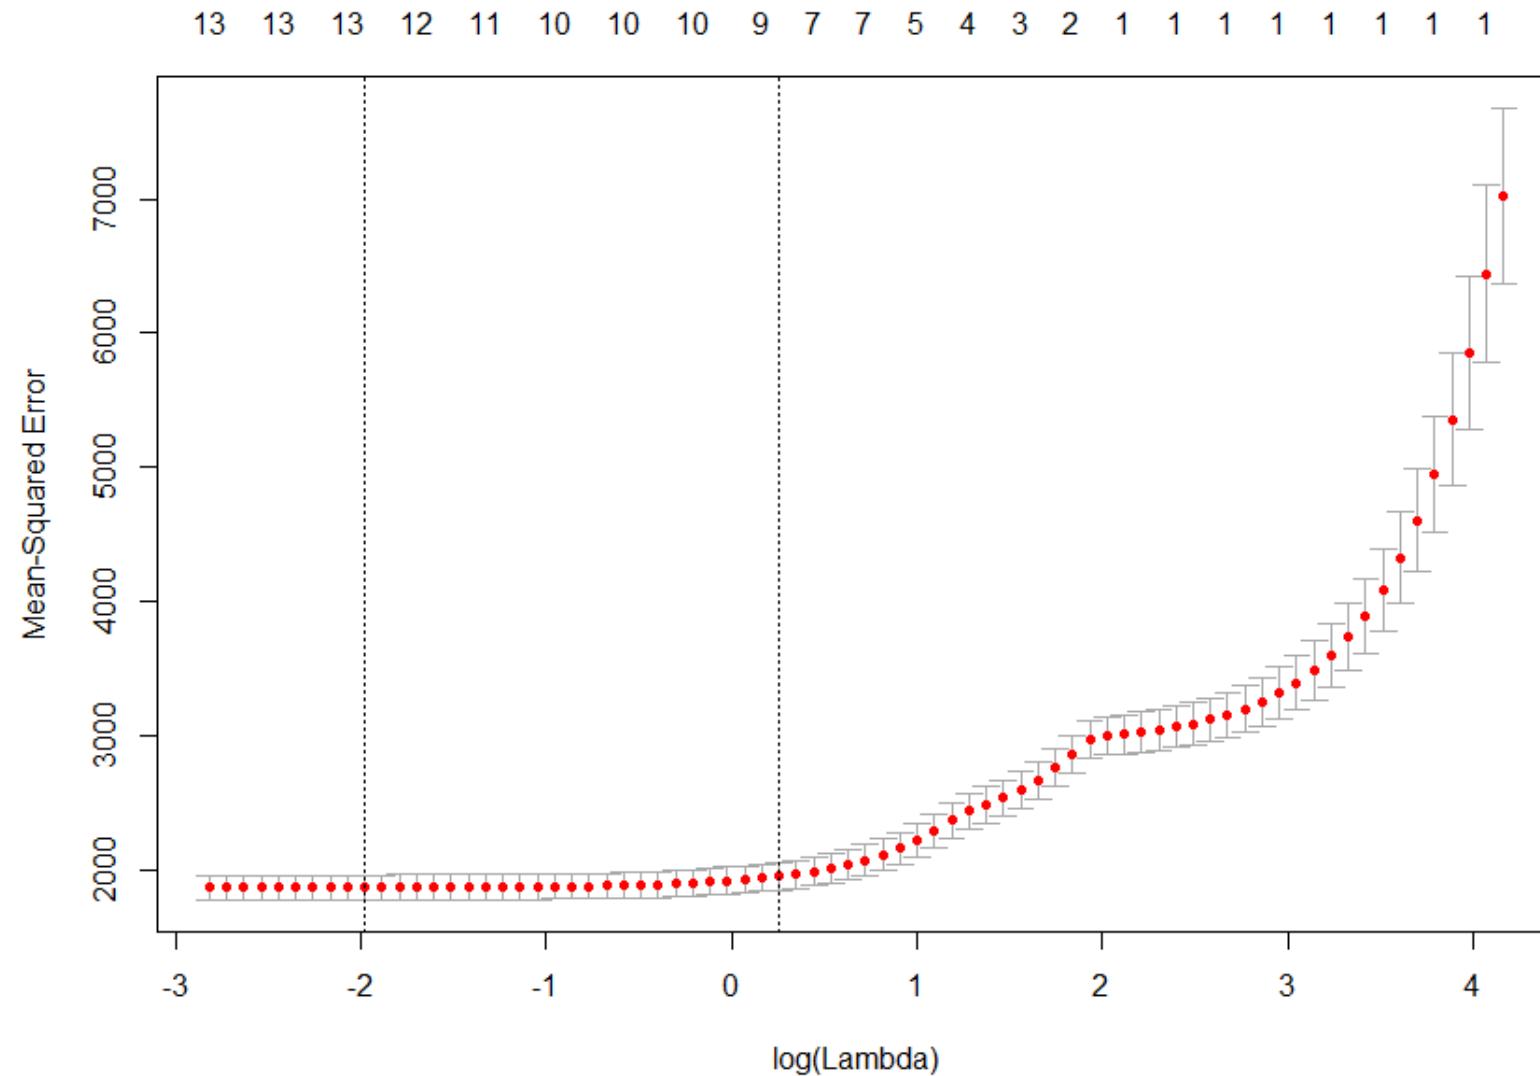
$$\min_{\beta} \left(\sum_{j=1}^N (y_j - \sum_{i=0}^d \beta_i \cdot x_i)^2 + \lambda \sum_{i=1}^d |\beta_i| \right)$$



CSE 7302c



LASSO: MSE



Ridge vs Lasso Regression

Ridge

- Reduces the magnitude of coefficients (does not make them zero)
- Solution is unique – however, doesn't do feature selection since it keeps all the terms
- Stable even if $n << p$
- Very stable even in multicollinear case

LASSO

- Values of irrelevant coefficients set to zero
- Good for feature selection
- Doesn't do well when $n << p$
- Out of sample performance can be impacted for multicollinear case

CS
7302C



Elastic-net

- A mix-between Ridge and LASSO regression

$$\min_{\beta} \left(\sum_{j=1}^N (y_j - \sum_{i=0}^d \beta_i \cdot x_i)^2 + \lambda \left(\sum_{i=1}^d (\alpha |\beta_i| + (1-\alpha) \beta_i^2) \right) \right)$$

- $\alpha=0$ is Ridge
- $\alpha=1$ is LASSO
- $\alpha=0.5$ is a equal mix of the two

- Keep in mind that Ridge regression can't zero out coefficients; thus, you either end up including all the coefficients in the model, or none of them.
- In contrast, the LASSO does both parameter shrinkage and variable selection automatically.
- If some of your covariates are highly correlated, you may want to look at the Elastic Net instead of the LASSO.

TIME SERIES FORECASTING

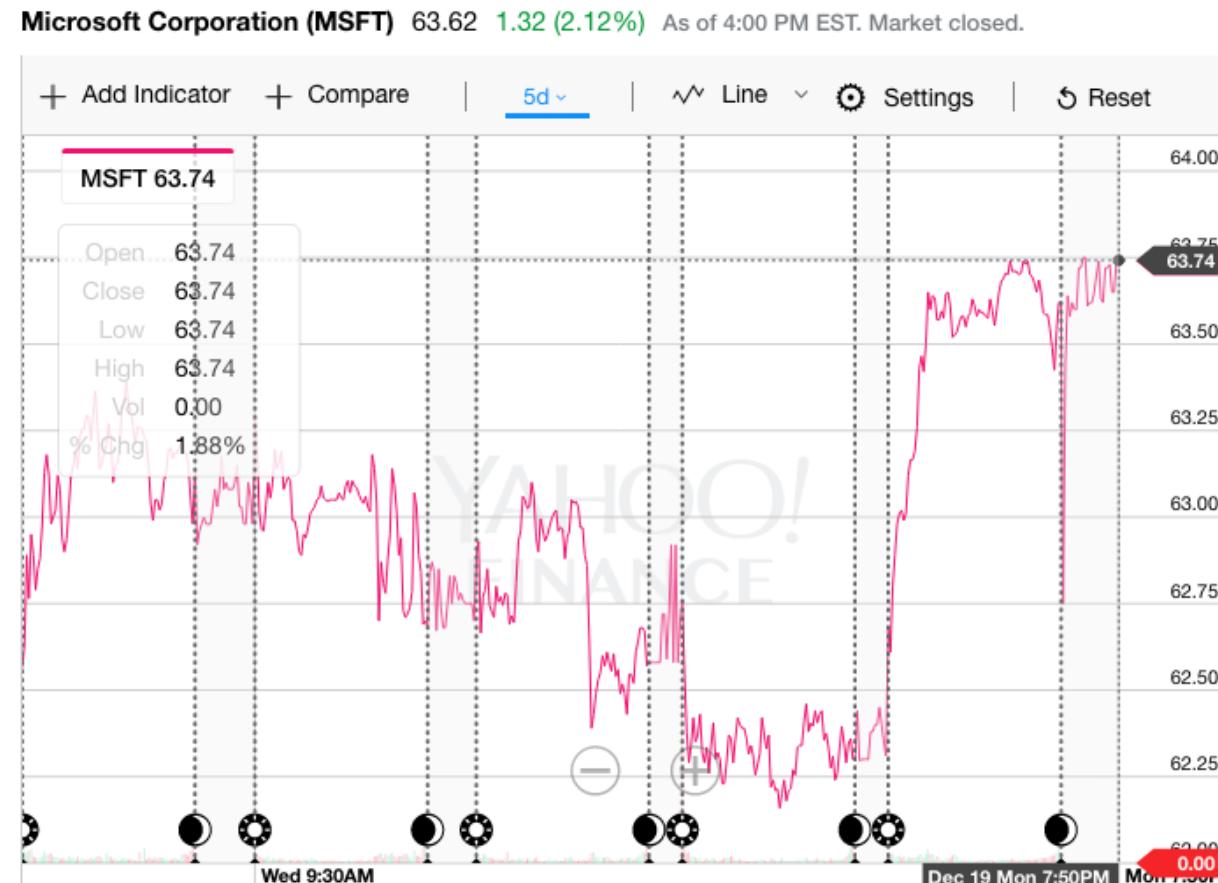
CSE 7302c



Why Time Series

Causal independent variables are

- Unknown to us
- Not available
- Might not fit the data well
- Difficult to forecast



Typical Time Series

$$\hat{y}_{t+1} = f(y_t, y_{t-1}, y_{t-2} \dots) + f(x_1, x_2, x_3 \dots)$$

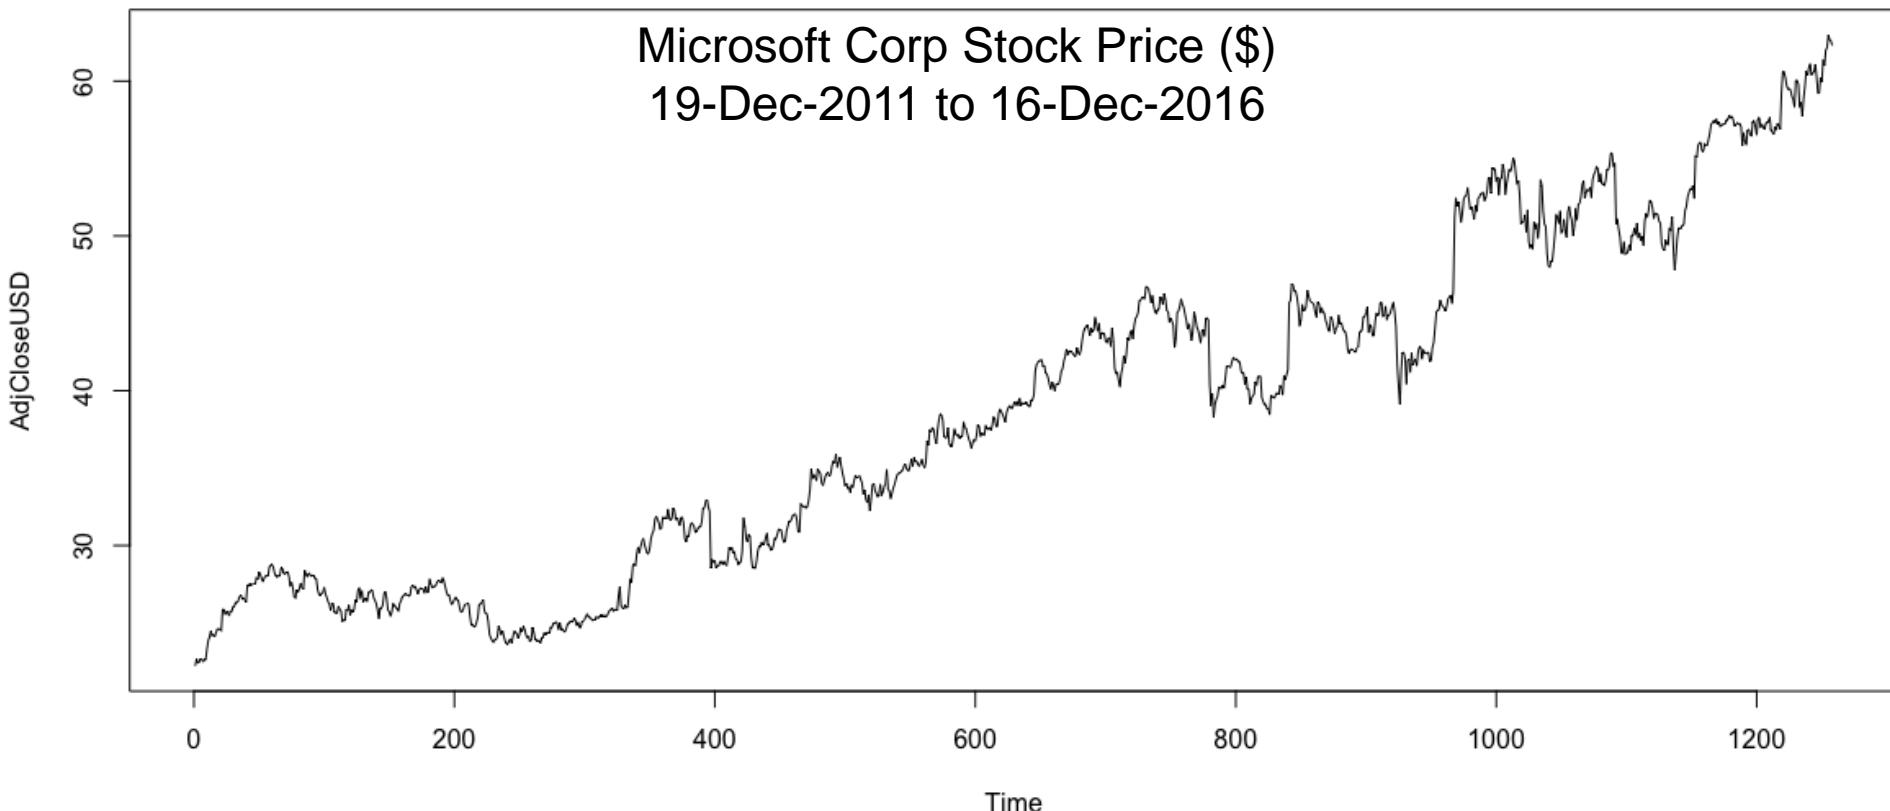
f can be linear or nonlinear

CSE 7302c



Components of Time Series

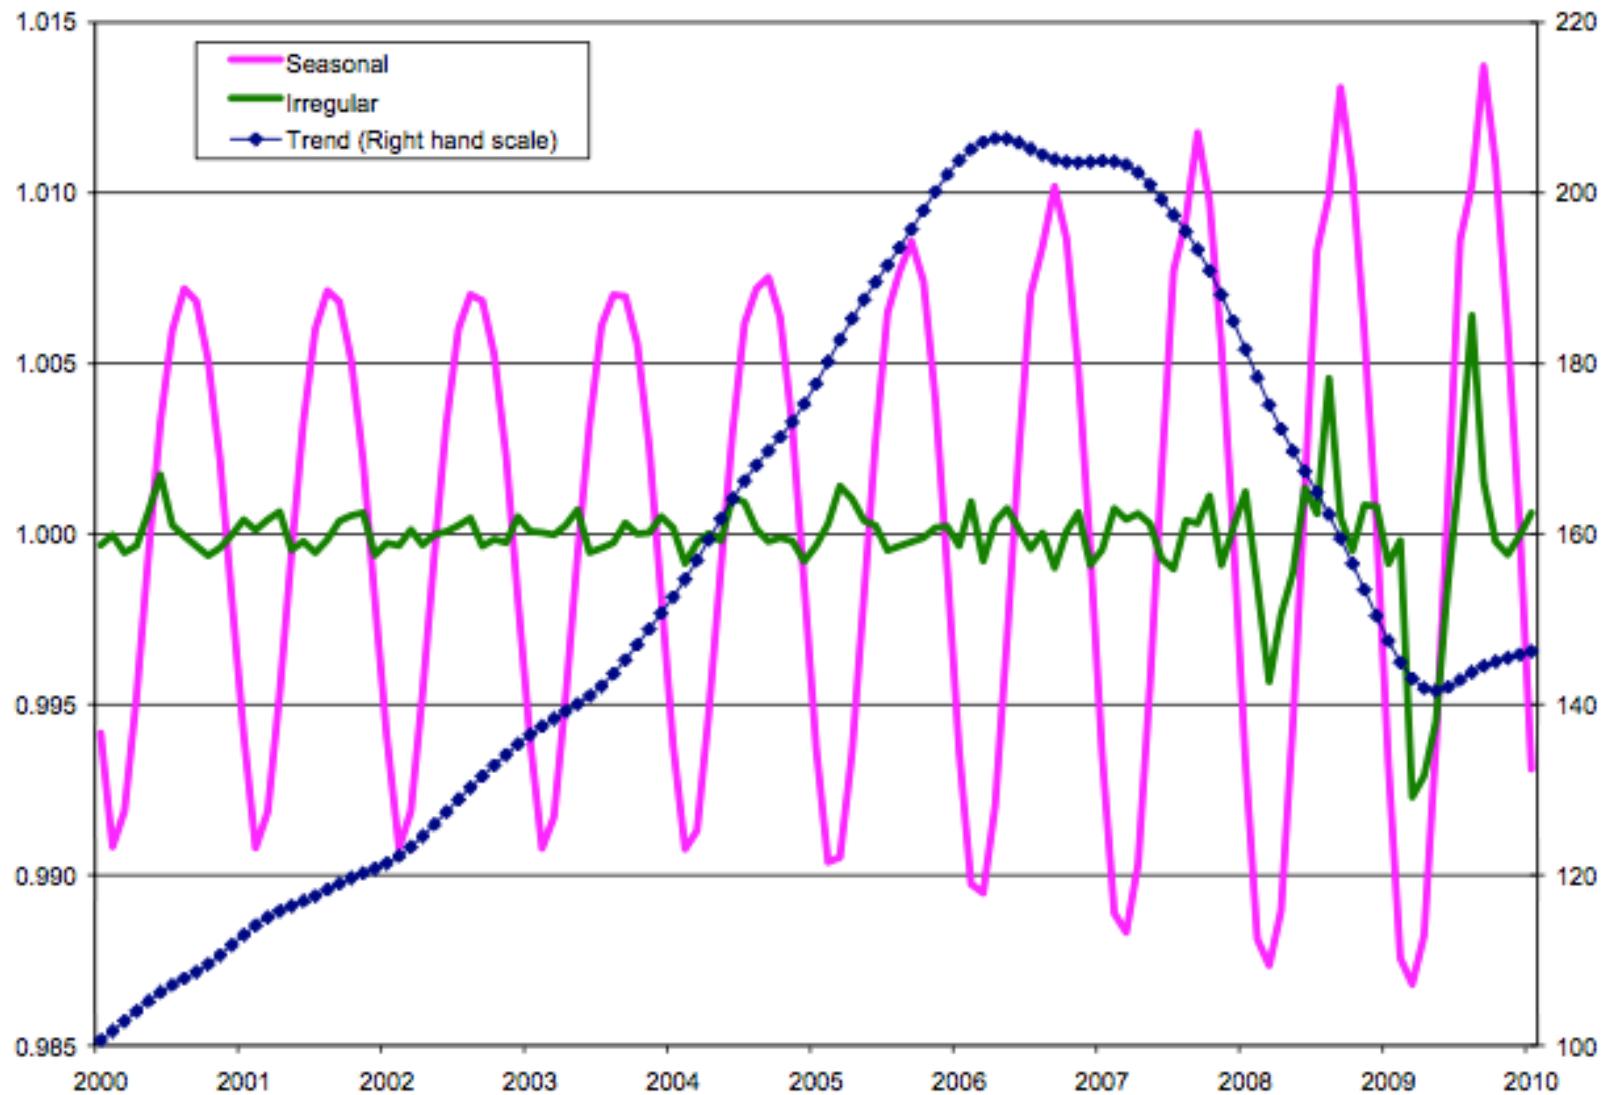
- Trend
- Seasonality
- Random component



CSE 7302C



Trend, Seasonality and Randomness



CSE 7302c



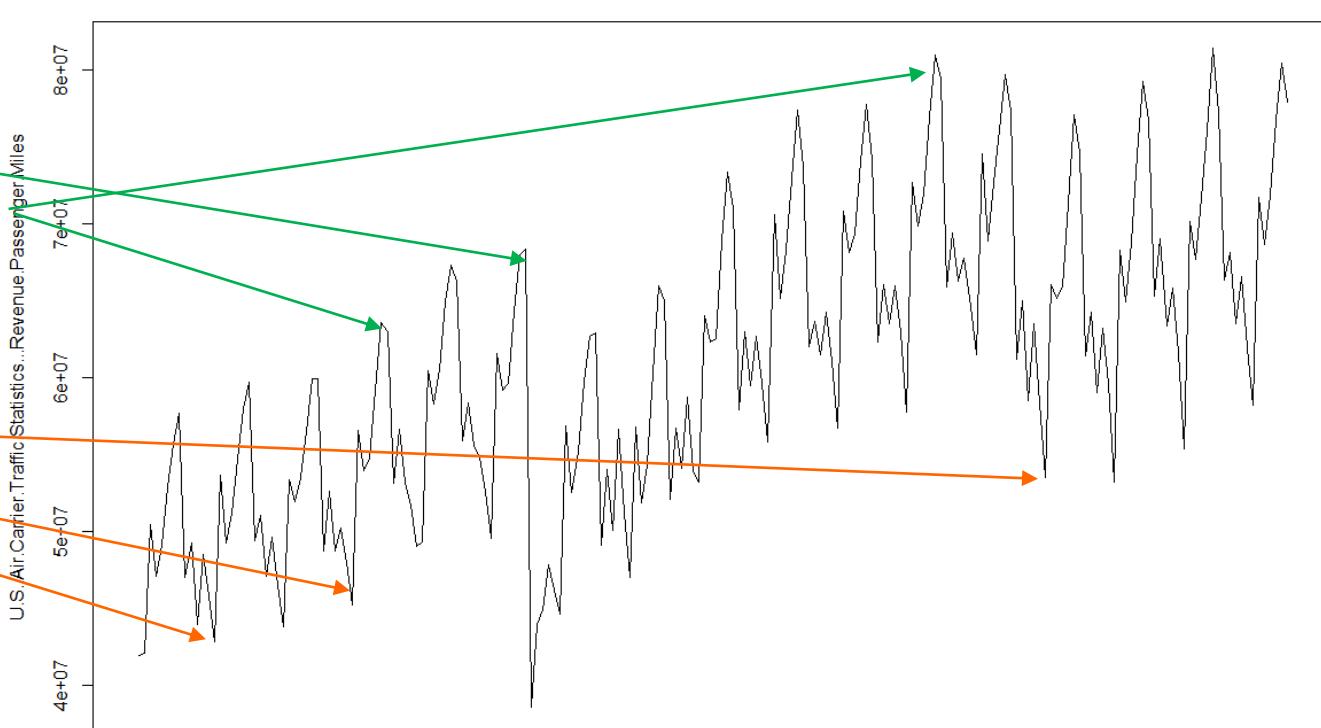
US Air Carrier Traffic – Revenue Passenger Miles ('000)

R code: `milestimeseries <- ts(miles, frequency = 12, start = c(1996,1))`

RPM

```
> milestimeseries <- ts(miles, frequency = 12, start = c(1996,1))
> milestimeseries
```

	Jan	Feb	Mar	Apr	May	Jun	Jul
1996	41972194	42054796	50443045	47112397	49118248	52880510	55664750
1997	45850623	42838949	53620994	49282817	51191842	54707221	57995025
1998	46514139	43769273	53361926	51968480	53515798	56460422	59939170
1999	47988560	45241211	56555731	53920855	54674958	59213000	63572248
2000	49045412	49306303	60443541	58286680	60533783	64903295	67346377
2001	52634354	49532578	61575055	59151645	59662416	64353323	67965298
2002	46224031	44615129	56897729	52542164	55116060	59745343	62664511
2003	51197175	47040806	56766580	51857453	54335598	60272900	65962215
2004	53979786	53179693	64035864	62340117	62530704	68866398	73335888
2005	59629608	55795165	70595861	65145552	68268899	72952959	77432998
2006	61035027	56729212	70799794	68120559	69352606	74099239	77798621
2007	63016013	57793832	72700241	69836156	71933109	76926452	80988340
2008	64667106	61504426	74575531	68906882	72725750	76162105	79707545
2009	58373786	53506580	66027341	65166300	65868254	71350227	77136799
2010	59651061	53240066	68307090	64953250	68850904	74474550	79304441
2011	61630362	55391206	70158268	67683558	71711448	76057910	81423215
2012	61940180	58243763	71696039	68669228	71887523	76760759	80499353
	Aug	Sep	Oct	Nov	Dec		
1996	57723208	47035464	49263120	43937074	48539606		
1997	59715433	49418190	51058879	47056048	49654209		



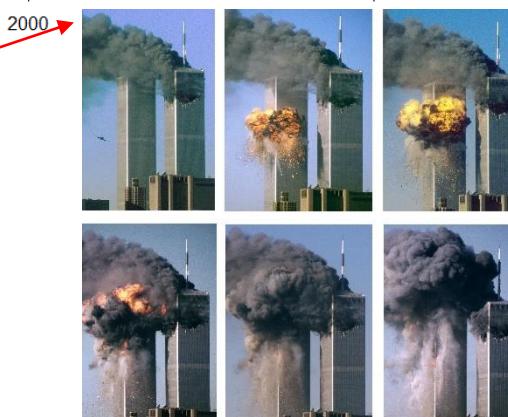
Data sources:

http://www.bts.gov/xml/air_traffic/src/index.xml

and <https://datamarket.com/data/set/281x/us-air-carrier-traffic-statistics-revenue-passenger-miles>

Last accessed: 31-Mar-2016

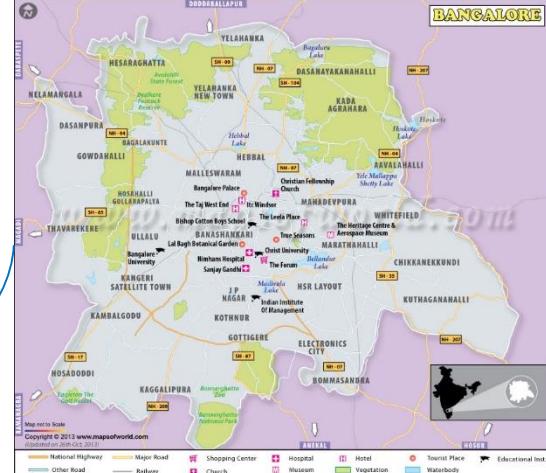
	Aug	Sep	Oct	Nov	Dec
1996	57723208	47035464	49263120	43937074	48539606
1997	59715433	49418190	51058879	47056048	49654209
1998	59927214	48751280	52578217	48734375	50208641
1999	63003663	53131972	56653901	53215500	51746821
2000	66256804	55900504	58373996	55590325	54822970
2001	68377080	38601868	43964788	44915764	47836501
2002	62944816	49096035	54019748	50106814	56656594
2003	64989766	52121480	56724551	54128776	58739845
2004	70961522	57881042	63021142	59453943	62680310



2010

7302C





HYDERABAD

2nd Floor, Jyothi Imperial, Vamsiram Builders, Old Mumbai Highway, Gachibowli, Hyderabad - 500 032
 +91-9701685511 (Individuals)
 +91-9618483483 (Corporates)

BENGALURU

Floors 1-3, L77, 15th Cross Road, 3A Main Road, Sector 6, HSR Layout, Bengaluru – 560 102
 +91-9502334561 (Individuals)
 +91-9502799088 (Corporates)

Social Media

- Web: <http://www.insofe.edu.in>
- Facebook: <https://www.facebook.com/insofe>
- Twitter: <https://twitter.com/Insofeedu>
- YouTube: <http://www.youtube.com/InsofeVideos>
- SlideShare: <http://www.slideshare.net/INSOFE>
- LinkedIn: <http://www.linkedin.com/company/international-school-of-engineering>

This presentation may contain references to findings of various reports available in the public domain. INSOFE makes no representation as to their accuracy or that the organization subscribes to those findings.